

Evaluation of the Performance of the Memory Tagging Extensions

Raphael Dichler

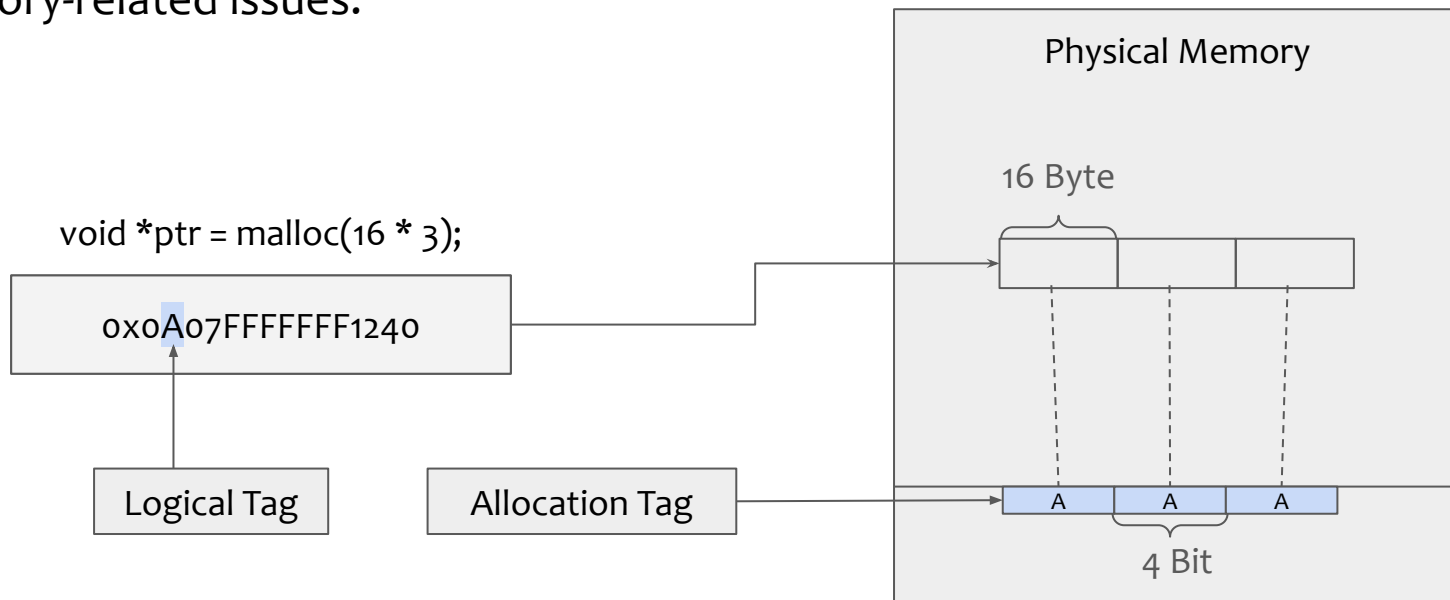
Supervisor: Martin Fink and Ilya Meignan--Masson



Memory Tagging Extension

In a nutshell

Is part of the ARMv8.5 specification, providing an approach to detect memory-related issues.



Outline



- ~~Memory Tagging Extension~~
- Motivation
- Experiments

Strong Security Guarantees

- MTE is a hardware feature designed to detect memory safety violations.

Performance Remains an Open Question

- Unknown impact on memory-bound and multi-threaded workloads.
- Integration in high-performance systems like databases is not well understood.

Key Challenges

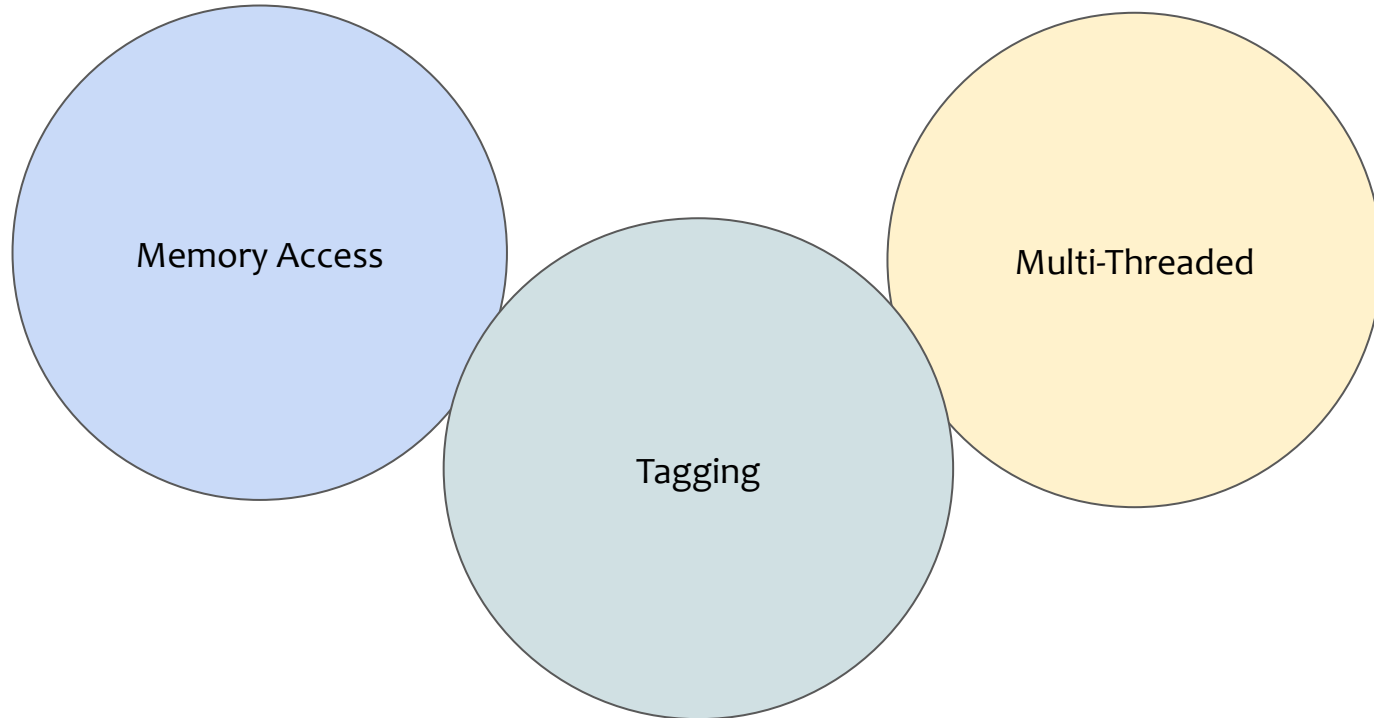
- Does tag fetching and checking introduce significant overhead?
- Extra steps in allocation and deallocation.
- Effects on caching, concurrency, and thread synchronization are unclear.

Outline



- ~~Memory Tagging Extension~~
- ~~Motivation~~
- Experiments

Key Areas of Performance Impact



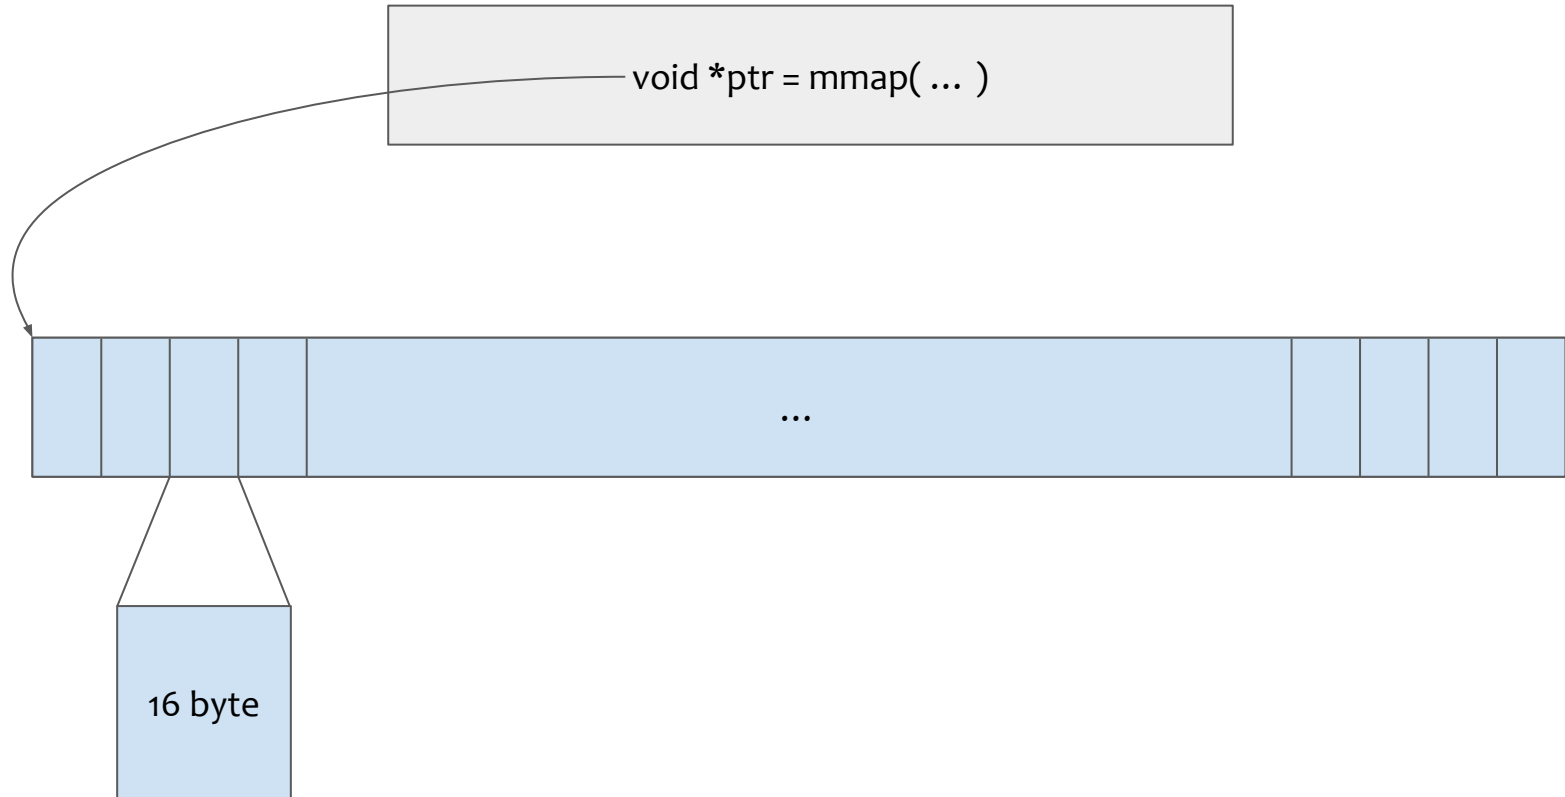
Idea

Analyze the overhead of non-contiguous memory access in the presence of MTE across different caches and main memory levels.


Experiment in a nutshell

- 1) Construct a linked list, which forms a loop
- 2) Measure the time to traverse over 100.000.000 nodes
- 3) Compare MTE-enabled with MTE-disabled


Experiment: Non-Contiguous Memory Access



Experiment: Non-Contiguous Memory Access



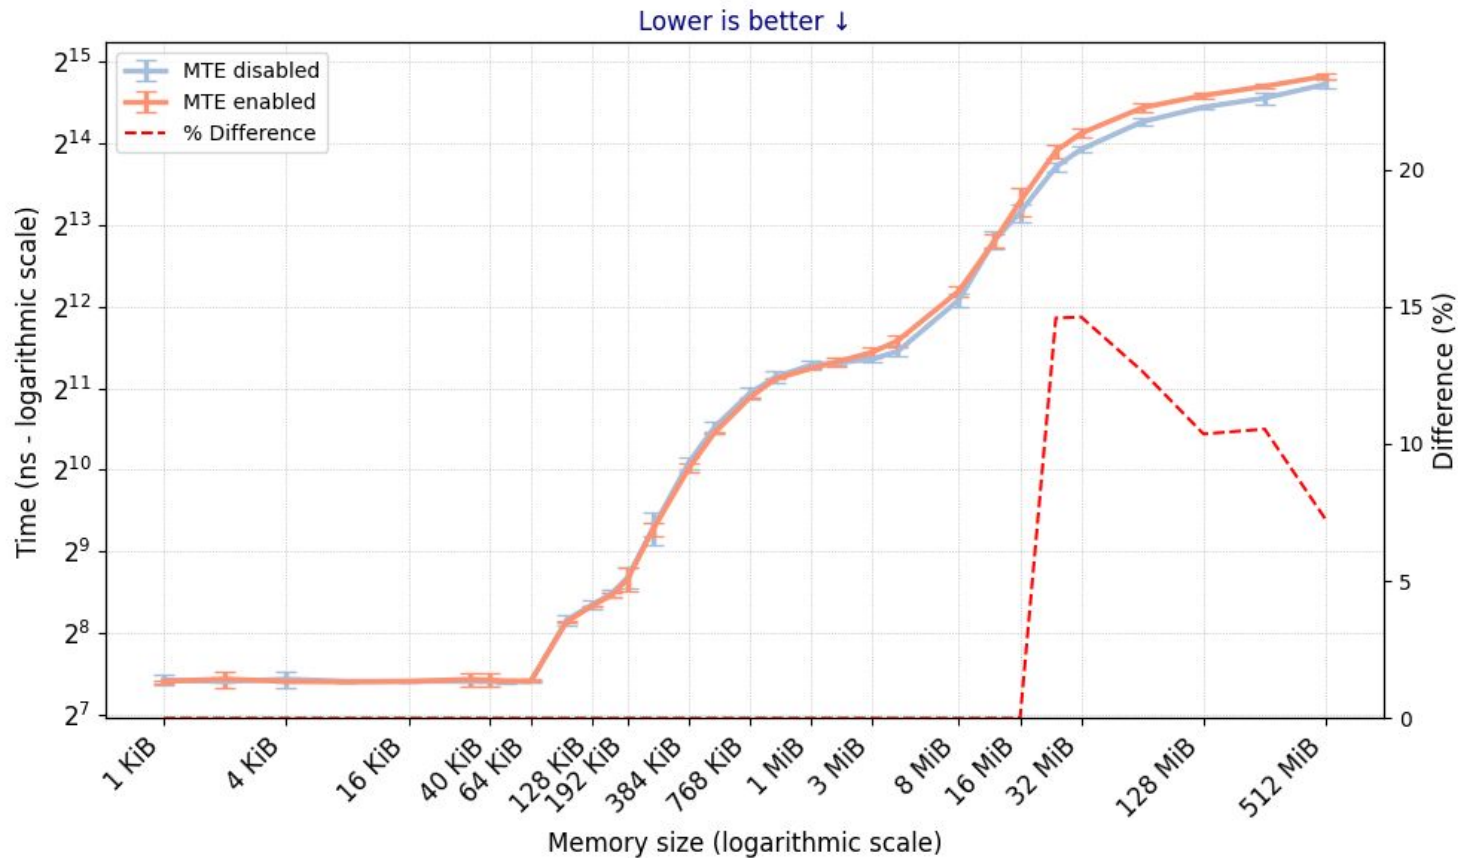
Rationale: Non-Contiguous Memory Access



- Data is small; all fits into L1 Cache
- As long the complete data fits in L1, we see no change in performance
- Holds for all Cache-Levels

* Due to specific cache replacement policies or access patterns, data may be evicted from L1/L2/L3 caches even if there appears to be sufficient space available.

Result: Non-Contiguous Memory Access



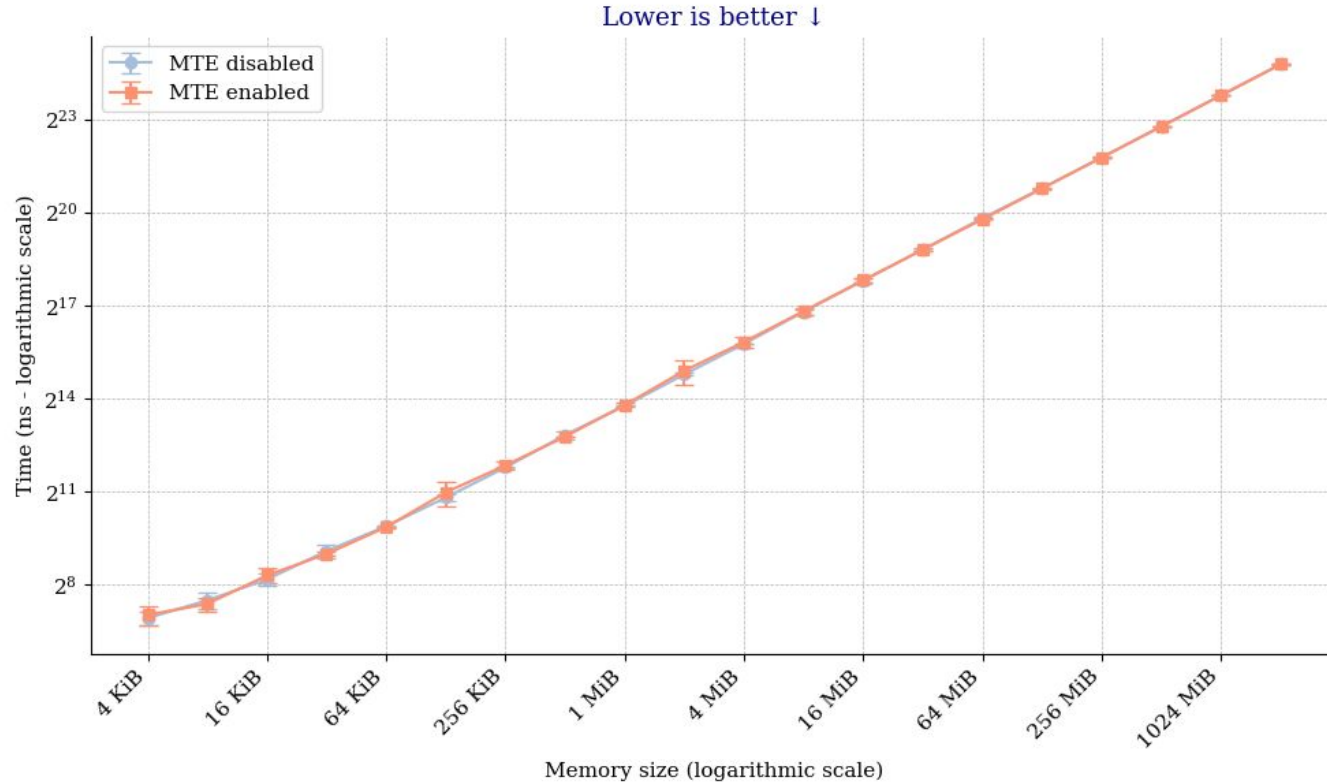
Idea

Analyze the overhead of memory access when accessing memory contiguously in the presence of MTE.

Experiment in a nutshell

- 1) Allocate a block of contiguous memory
- 2) Measure the time to iterate once
- 3) Compare MTE-enabled with MTE-disabled

Result: Contiguous Memory Access



Experiment: Tagging

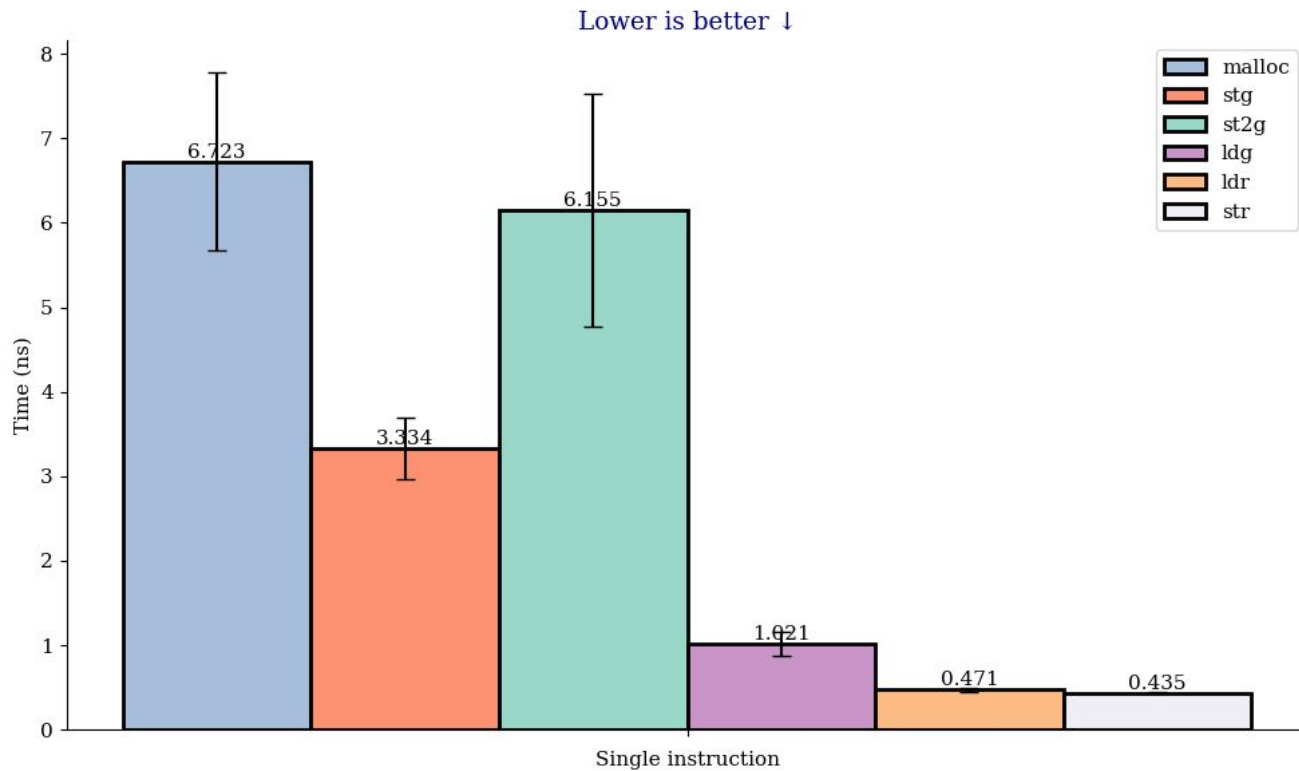
Idea

Compare the execution time to tag memory with load and store.

Experiment in a nutshell

- 1) Allocate a block of contiguous memory
- 2) Measure the time of tagging (stg, st2g), loading, or storing data

Result: Tagging



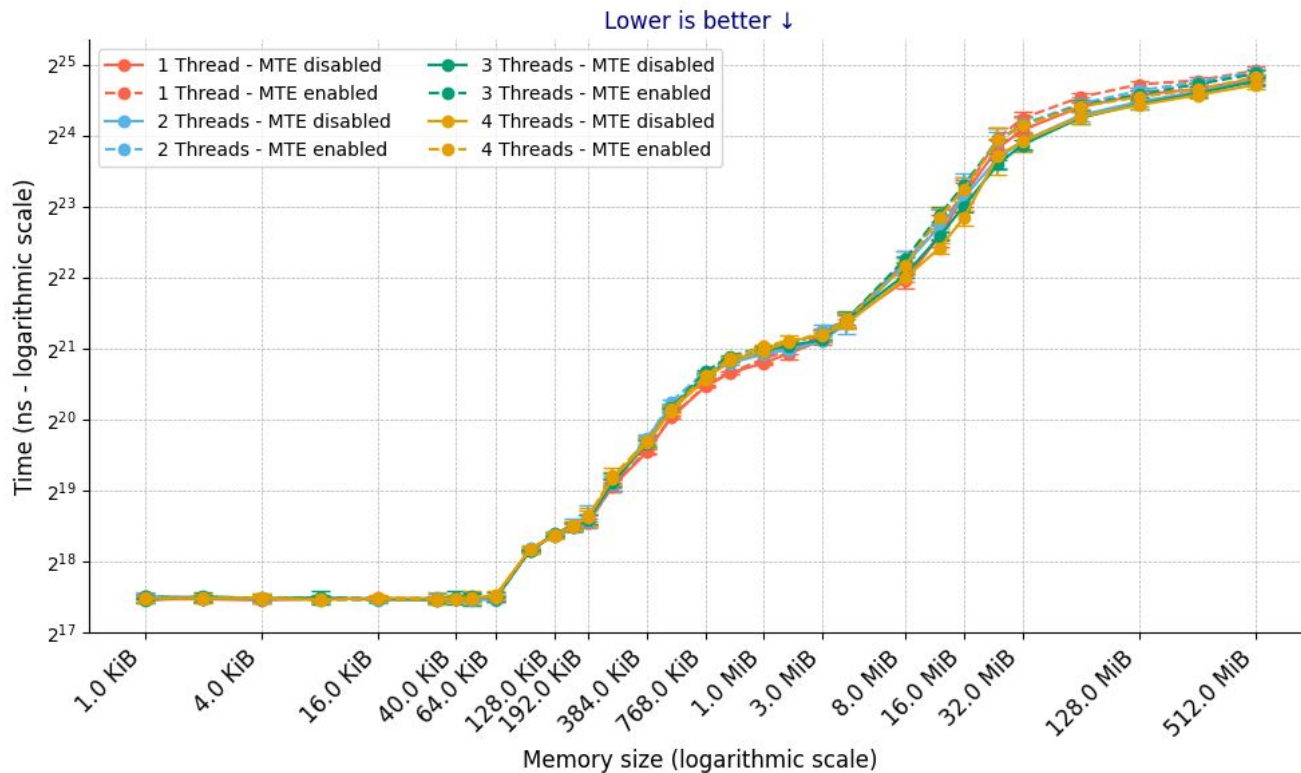
Idea

Analyze the overhead of concurrent read and write when accessing memory in the presence of MTE.


Experiment in a nutshell

- 1) Construct a linked list like in the non-contiguous memory access experiment.
- 2) Create multiple threads and pass the start of the linked list to them.
- 3) Measure the time to traverse over 100.000.000 nodes by each thread

Result: Concurrent Read



Result: Concurrent Write



Outline



- ~~Memory Tagging Extension~~
- ~~Motivation~~
- ~~Experiments~~

Key Areas of Performance Impact

As long data resides inside the **cache or is prefetchable** MTE introduces **no overhead**

If data is **not prefetchable nor located inside** the cache MTE introduces **5-15% overhead**

Tagging takes about **x7 times longer** than load/store operations

Concurrent reads don't introduce any overhead when MTE is enabled

Concurrent writes introduce an overhead of 5-15% when MTE is enabled

Backup

Idea

Analyze the overhead of Compare-And-Swap (CAS) when accessing memory in the presence of MTE.

Experiment in a Nutshell

- 1) Allocate a block of shared memory.
- 2) Create multiple threads and pass them the pointer to the memory.
- 3) Each thread increments the shared value by one 1.000.000 times.
- 4) Compare MTE-enabled with MTE-disabled

Result: Tagging

