

Design and Implementation of a Binary Translator from AArch64 to a Custom Intermediate Representation

Konstantin Garbers

Advisor: Martin Fink

Chair of Computer Systems

<https://dse.in.tum.de/>



29.10.2024 – 28.02.2025

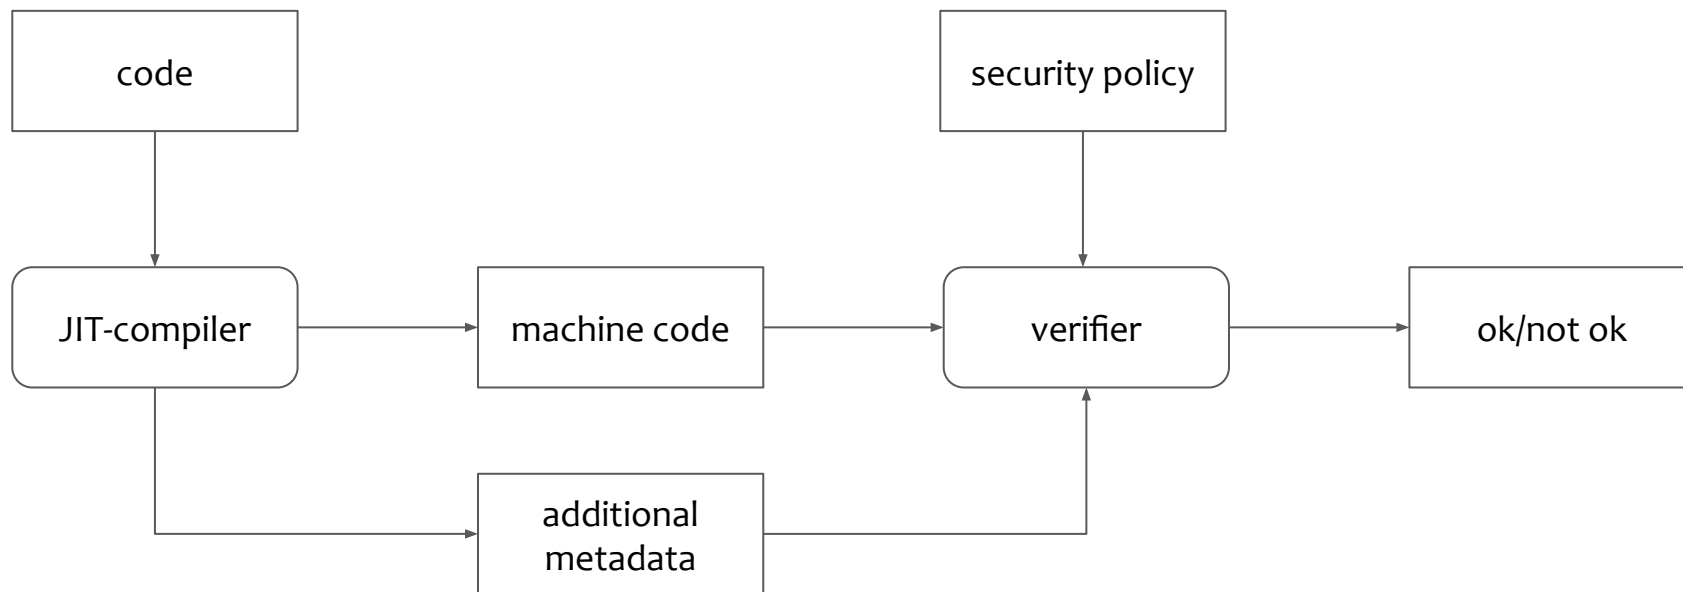
Introduction: JIT-Compilers



- Code execution tools used in the JVM or in browsers
- Interpretes code and compiles certain parts of code to machine code
- Compiled machine code is source of security issues

Introduction: TrustNoJit (TNJ)

Library that uses the **proof-carrying-code** framework to verify JIT-compiled code



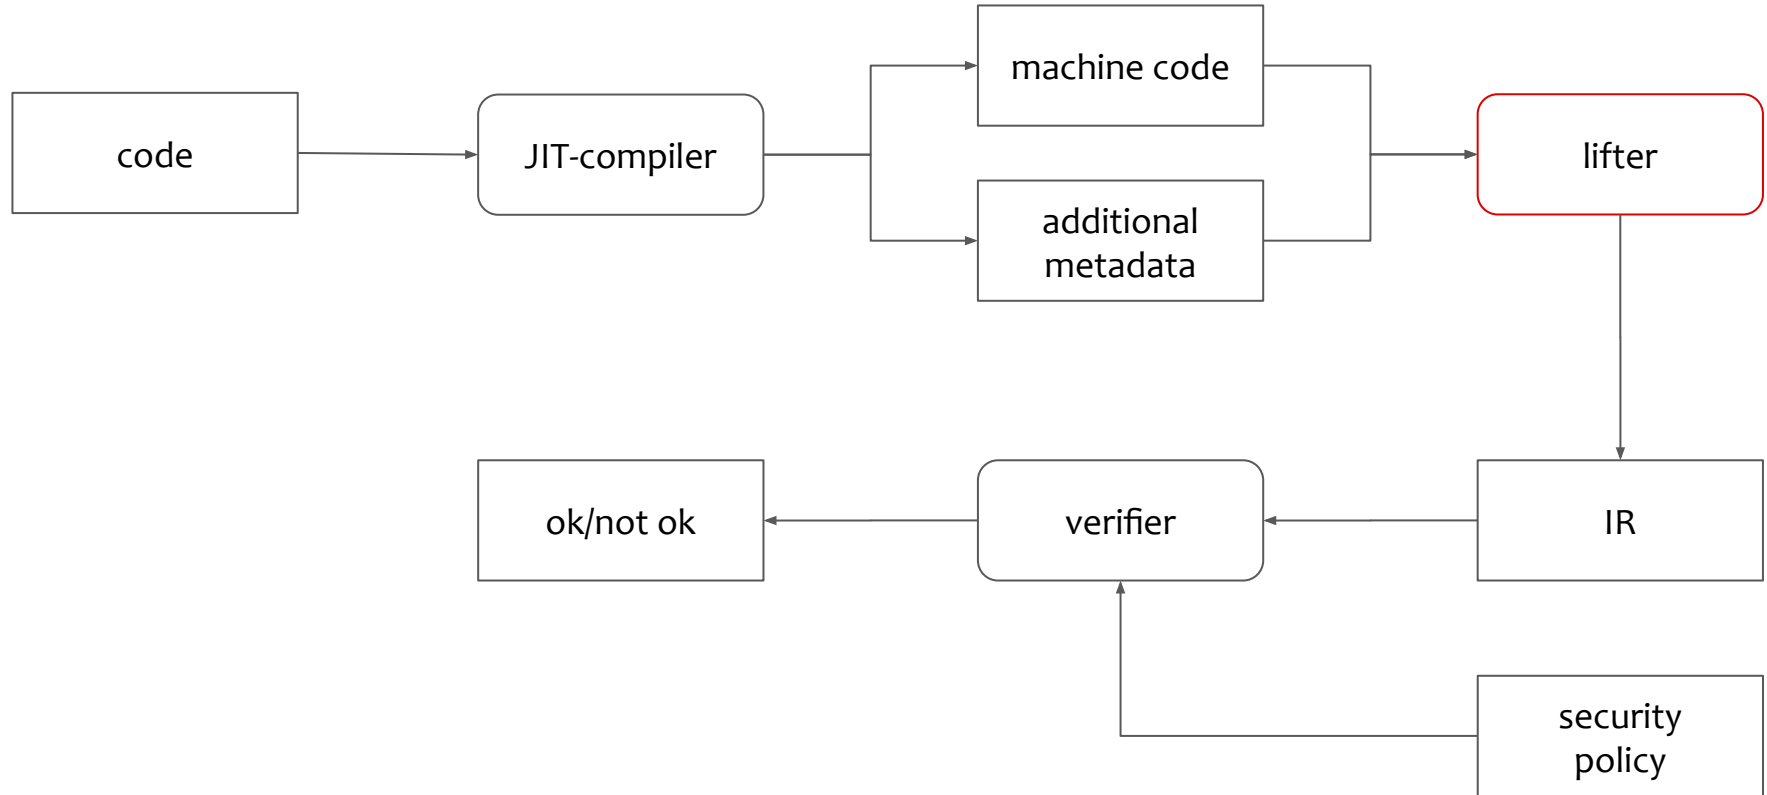
Problem: Machine code is hard to verify!

- **Difficult control flow** tracking
- Verification requires intricate knowledge of **side effects** of assembly instructions
- Verifier needs to be rewritten for every language

Solution:

- Transform assembly code into an Intermediate Representation (IR)

Solution: Integration of a lifter into TNJ



Outline



- ~~Motivation~~
- Background
- Design
- Implementation
- Evaluation

Background: AArch64

64-bit execution state of ARMv8-A

- Every instruction is 4 bytes

Challenges:

- 5 addressing modes
- 32 general-purpose registers
- 4 condition flags
- More than 300 base instructions and multiple extensions

Background: Assembly IR (AIR)

```
0x00: ldr x0, [x1, #0x80]
```

BasicBlock (BB)

```
block_0:  
v0 = i64.read_reg "x1"  
v1 = i64.add v0, 0x80  
v2 = i64.load v1  
i64.write_reg v2, "x0"
```

AArch64 Code

AIR

Outline

- ~~Motivation~~
- ~~Background~~
- Design
- Implementation
- Evaluation

Problem: Lifter Can Not Translate Backward Jumps

AArch64 Code

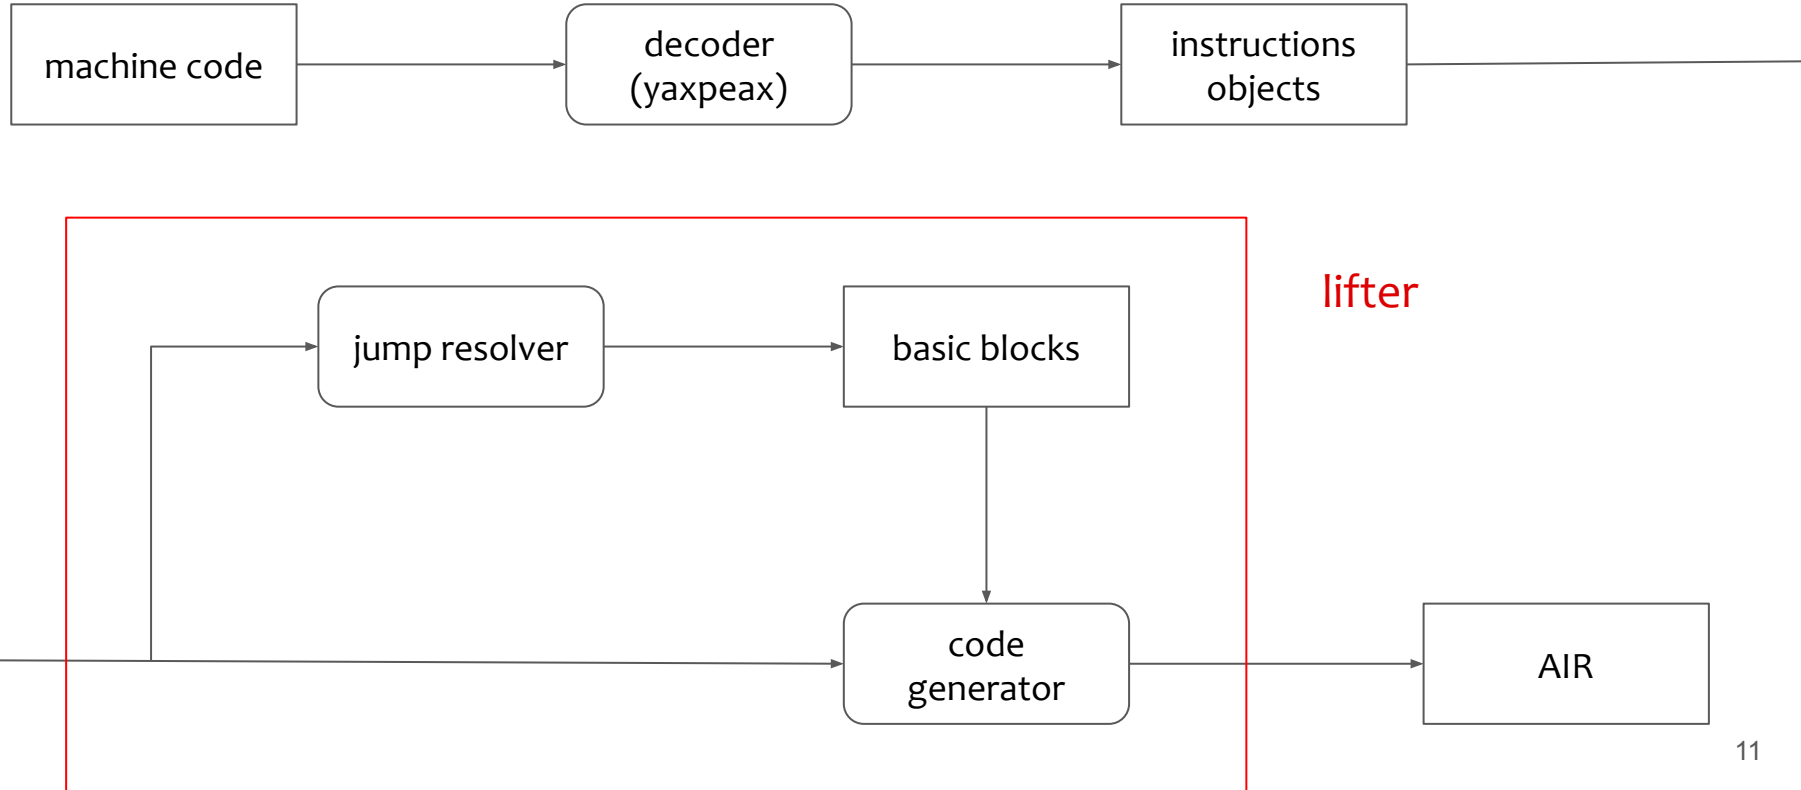
0x00: add x1, x2, x3

0x04: b 0x08

0x08: b 0x04

- AIR consists of BasicBlocks (BB)
- Jumps **jump to the beginning** of BBs
- Jumps may only be **at the end** of a BB

Lifter Workflow



Solution: Create BBs in A Separate Translation Pass

AArch64 Code

```
0x00: add x1, x2, x3
0x04: b 0x08
0x08: b 0x04
```

AIR

```
block_0:
  v0 = i64.read_reg "x2"
  v1 = i64.read_reg "x3"
  v2 = i64.add v1, v0
  i64.write_reg v2, "x1"
  jump block_4
```


```
block_4:
  jump block_8
```

```
block_8:
  jump block_4
```

Unsupported Instructions

Mark destination registers of unsupported instructions as opaque

```
umov x0, v1.s[0]
```




```
v0 = i64.opaque  
i64.write_reg v1, "x0"
```

Represent unknown value
using opaque value

System calls, hypervisor calls or function calls modify registers untraceably

0x00: blr x0



```
v0 = i64.read_reg "x0"  
dynamic_jump v0  
invalidate_regs
```

Outline



- ~~Motivation~~
- ~~Background~~
- ~~Design~~
- Implementation
- Evaluation

Implementation

- Developed in **rustc 1.84.0-nightly**
- Uses the **yaxpeax¹** library to decode machine code
- Supports **129 out of 301** base variants

¹yaxpeax: <https://github.com/iximeow/yaxpeax-x86>

Outline

- ~~Motivation~~
- ~~Background~~
- ~~Design~~
- ~~Implementation~~
- Evaluation

Evaluation: Setup

Hardware:

- AMD Ryzen 5 5500U CPU
- 16 GiB DRAM
- Ubuntu 22.04.05 LTS

Test Binaries:

- WASM-Applications from the Sightglass²-benchmarking suite
- Compilation to AArch64 binaries using standard compilation settings

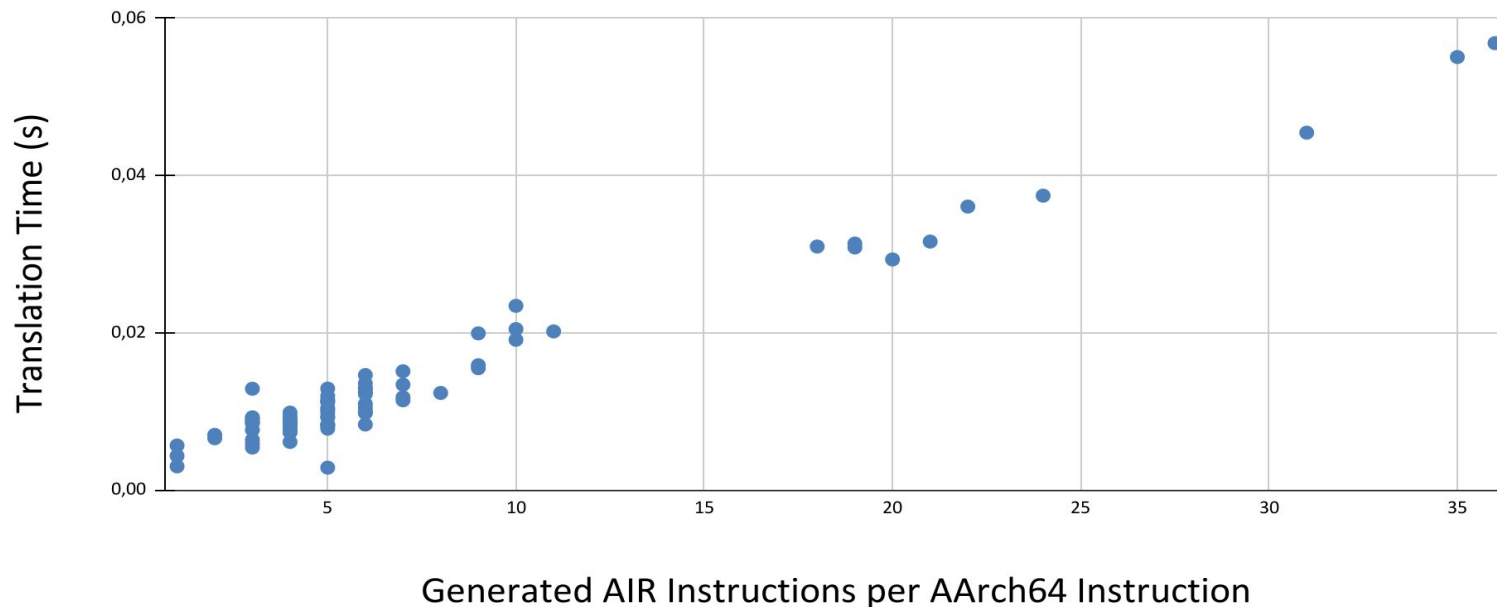
²Sightglass: <https://github.com/bytecodealliance/sightglass>

Evaluation: Research Questions

1. Is lifter **fast** enough for real-world-settings?
2. Is the lifter **usable** in real-world settings?


Insight: Translation Speed Correlates To Generated AIR-Instructions

Comparison of Instruction Translation Time to Generated AIR Instruction



Insight: Lifting Time Varies a Lot Across Different Applications

Comparison of Objdump to Lifter Execution Time



Lower Values -> Better

Evaluation: Research Questions

1. Is lifter **fast** enough for real-word-settings?
 - A single instruction takes **20 μ s-80 μ s** to lift
 - Average function will take **2ms-11ms** to lift

2. Is the lifter **usable** in real-world settings?
 - Machine decoding dependency **yaxpeax is a current limitation**
 - Supports on average more than **99.7%** of all tested instructions
 - Supports more than **99%** of instructions in every tested application

Related Works: State-of-the-art

Verifier Frameworks that lift AArch64 into an IR:

- angr¹: Python framework that lifts AArch64 into a VEX IR
- Miasm²: Python framework that lifts AArch64 into a custom IR
- BAP³: OCaml framework that lifts AArch64 into a custom IR

¹ angr: <https://github.com/angr/angr>

² Miasm: <https://github.com/cea-sec/miasm>

³ BAP: <https://github.com/BinaryAnalysisPlatform/bap>

Related Works: A Custom Lifter is Advantageous



Integration with TNJ

- **Smooth data flow** between the verifier and lifter
- **Direct access** to TNJ's built-in functionalities

Execution Speed

- **Single-purpose** application
- **Lower-level programming language** for faster execution time

Tailored Intermediate Representation (IR)

- IR is specifically designed for **TNJ's requirements**

Contributions

- Design of a custom intermediate language AIR
- Implementation of an AArch64 to AIR lifter

Evaluation

- Lifter supports more than **99.7%** of all encountered instructions
- Decoding dependency yaxpeax is a **bottleneck**
- Future aim is support of **function calls, pointer authentication and memory ordering operations**

Backup

Implementing Flag Registers in AIR

Represent flags using **separate** registers

Reading Flags:

```
v0 = i1.read_reg "z"
```

Writing Flags:

```
i1.write_reg v0, "z"
```

Represent flags using a **single** registers

Writing immediates:


```
i8.write_reg 0xf, "flags"
```

Implementing Register Sizes in AIR

Supported Sizes: **BOOL, l8, l16, l32, l64, l128**

Edge case requires support for l128:

```
smulh x1, x2, x3
```



```
v0 = i64.read_reg "x2"  
v1 = i64.read_reg "x3"  
v2 = i64.imul v0, v1  
v3 = i128.ashr v2, 0x40  
i64.write_reg v3, "x1"
```

Implementing Flags in AIR

Represent flags using **separate** registers

Reading Flags:

```
v0 = i1.read_reg "z"
```

Writing Flags:

```
i1.write_reg v0, "z"
```

Generated instructions per application

