

# Microarchitectural Analysis of CHERI on the Morello Platform

Yude Jiang

Advisor: Martin Fink, Dr. Masanori Misono


Chair of Computer Systems

<https://dse.in.tum.de/>




28.11.2024 – 28.03.2025

# Memory safety bugs are a critical problem in Software



[1] Chromium high-severity security bugs



[2] MSRC Microsoft CVEs

Possible Solution:  
**Hardware-based memory protection mechanisms**

[1] <https://www.chromium.org/Home/chromium-security/memory-safety>

[2] <https://msrc.microsoft.com/blog/2019/07/a-proactive-approach-to-more-secure-code>

## Capability Hardware Enhanced RISC Instructions

- **Capabilities:**

Unforgeable, permission-tagged, and bounds-checked references that replace raw pointers

- **Memory Safety Enforcement:**



Hardware checks on capability bounds and validity prevent spatial and temporal memory vulnerabilities

- **Execution Models:**

- **Hybrid:** Legacy and capability-aware code coexist
- **Purecap:** All pointers are capabilities - full fine-grained memory safety

# Background: Morello

- Research Platform provided by Arm
- Extends the Armv8.2 ISA with CHERI
- Uses a merged register file for capability registers
- Hardware prototype based on Neoverse N1 shipped in 2022



[1] Morello SOC on board

## **Problem:**

Limited performance data on CHERI/Morello hinders accurate evaluation of its security-performance tradeoff.

## **Solution:**

**Perform microbenchmark-based instruction analysis to expose latency and throughput characteristics that inform optimization.**

# Research Objectives

**RO1:** What is the effect of Capabilities on memory instructions?

**RO2:** Does bounds checking introduce overhead?

**RO3:** What are the Latency and Throughput of Morello instructions?

We measure the following metrics for a given Instruction:

- **Latency:** How long an instruction takes to execute
  - Measured by introducing dependencies between instructions
  - Indicates how long an instruction stalls dependent operations
- **Throughput:** How many instructions can be executed in one cycle
  - Measured by eliminating dependencies between instructions
  - Indicates how effectively the CPU can pipeline and overlap executions

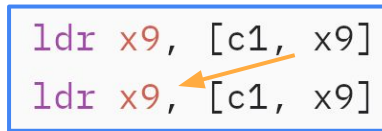


Diagram showing two instructions in a blue box: `ldr x9, [c1, x9]` and `ldr x9, [c1, x9]`. An orange arrow points from the `x9` register in the second instruction back to the `x9` register in the first instruction, indicating a data dependency.

Dependent instructions

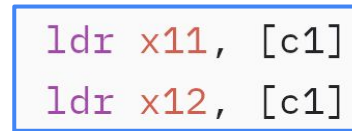


Diagram showing two instructions in a blue box: `ldr x11, [c1]` and `ldr x12, [c1]`. There is no dependency between these instructions as they access different registers.

Independent instructions

# Memory Instructions

Nocap results (Baseline)

Instruction	Throughput (i/c)	Latency (c/i)
LDR	1.986	4.010
STR	1.972	-
LDP	0.996	4.008
STP	0.994	-
LDAR	1.990	-
STLR	1.985	-

Purecap results

Instruction	Throughput (i/c)	Latency (c/i)
LDR	1.990	4.009
STR	1.981	-
LDP	0.996	-
STP	0.996	-
LDAR	1.993	-
STLR	1.990	-
LDR_CAP	1.990	4.005
LDP_CAP	0.995	4.009
STR_CAP	0.997	-
STP_CAP	0.499	-

No performance difference between  
memory addressing modes



# Memory Instructions

Nocap results (Baseline)

Instruction	Throughput (i/c)	Latency (c/i)
LDR	1.986	4.010
STR	1.972	-
LDP	0.996	4.008
STP	0.994	-
LDAR	1.990	-
STLR	1.985	-

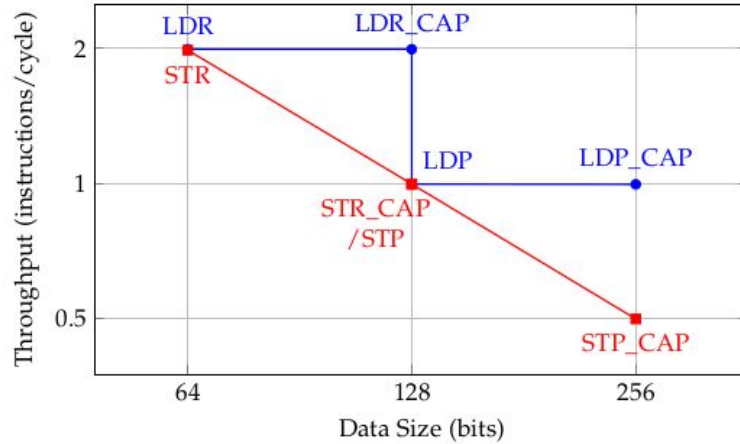
Purecap results

Instruction	Throughput (i/c)	Latency (c/i)
LDR	1.990	4.009
STR	1.981	-
LDP	0.996	-
STP	0.996	-
LDAR	1.993	-
STLR	1.990	-
LDR_CAP	1.990	4.005
LDP_CAP	0.995	4.009
STR_CAP	0.997	-
STP_CAP	0.499	-

No performance difference between memory addressing modes

No performance penalty for instructions enforcing access order

# Memory Instructions



Instruction throughput vs data size

Instruction	Offset	In-place
STR	1.981	→ 1.984
STP	0.996	→ 0.333
STR_CAP	0.997	→ 0.333
STP_CAP	0.499	→ 0.166

Offset/in-place store throughput comparison

Load instructions scale based on if they are single or paired

Store instructions scale linearly with the stored datasize

Store throughput decreases 3x when performing in-place stores sized over 64 bits

# Capability Instructions

Instruction	Throughput (i/c)	Latency (c/i)
CVT_TOCAP	0.998	2.006
CVT_TOPTR	0.997	2.007
CVTP_TOCAP	0.997	2.004
CVTP_TOPTR	0.998	2.005
CVTD_TOCAP	0.998	2.006
CVTD_TOPTR	0.997	1.004
CFHI	2.800	1.005
CTHI	2.818	1.007

Most conversion instruction  
perform the same

# Capability Instructions

Instruction	Throughput (i/c)	Latency (c/i)
CVT_TOCAP	0.998	2.006
CVT_TOPTR	0.997	2.007
CVTP_TOCAP	0.997	2.004
CVTP_TOPTR	0.998	2.005
CVTD_TOCAP	0.998	2.006
CVTD_TOPTR	0.997	1.004
CFHI	2.800	1.005
CTHI	2.818	1.007

Most conversion instruction  
perform the same

CVTD\_TOPTR has half the  
latency of the others

# Capability Instructions

Instruction	Throughput (i/c)	Latency (c/i)
CVT_TOCAP	0.998	2.006
CVT_TOPTR	0.997	2.007
CVTP_TOCAP	0.997	2.004
CVTP_TOPTR	0.998	2.005
CVTD_TOCAP	0.998	2.006
CVTD_TOPTR	0.997	1.004
CFHI	2.800	1.005
CTHI	2.818	1.007

Most conversion instruction  
perform the same

CVTD\_TOPTR has half the  
latency of the others

CFHI and CTHI are highly  
performant

## **RO1: What is the effect of Capabilities on memory instructions?**

- Loading capabilities is just as performant as loading integer addresses
- Storing capabilities is less performant due to store buffer limitations

## **RO2: Does bounds checking introduce overhead?**

- No performance difference between addressing modes

## **RO3: What are the Latency and Throughput of Morello instructions?**

- Most conversion instructions have high latency of around 2 cycles/instruction
- Throughput of 1 instruction/cycle indicates effective pipelining or multiple execution units

- **What did we do?**

We performed microbenchmark based analysis of the Morello CPU to expose its Latency and Throughput characteristics.

- **What did we find?**

- Capability loads are efficient
- Stores and conversions can introduce overhead
- Bounds checking is inexpensive


- **What could be done in the future?**

- Methodology extension
- Application level analysis

# Backup



# Background: CHERI Capabilities



- Integer Address:** Conventional 64-bit integer address
- Bounds:** Compressed addresses which describe accessible memory area
- Object Type:** Indicates if and how a capability is sealed
- Permissions:** Indicates how a capability can be used
- Validity Tag:** Tracks if the capability is valid

# Background: CHERI Software Modes

	Pure-capability Mode	Hybrid Mode
Memory access via capabilities	✓	✓
Memory access via integer addresses	✗	✓ *
Fine-grained memory safety	✓	✗
Compatibility with legacy code	✗	✓

\* Bound/Permission checks are performed against the Default Data Capability (DDC)


# Background: CHERI Architectural Rules

## Provenance Validity:

A valid capability can only be derived from another valid capability.

## Capability monotonicity:

A capability cannot have greater permissions or bounds than the capability it was derived from.



```
ldr_latency_loop:
    mov x9, #0
    ldrl_loop:
        cbz x0, ldrl_end
        sub x0, x0, #1
        ...
        ldr x9, [c1, x9]
        ldr x9, [c1, x9]
        ...
        b ldrl_loop

    ldrl_end:
        ret
```

```
ldr_throughput_loop:
    cbz x0, ldrt_end
    sub x0, x0, #1
    ...
    ldr x11, [c1]
    ldr x12, [c1]
    ...
    b ldr_throughput_loop

    ldrt_end:
        ret
```