

# Online OS reconfiguration for Cloud DBMS

Hristina Yordanova Grigorova

Advisor: Ilya Meignan--Masson

Systems Research Group

<https://dse.in.tum.de/>



28.04.2025 – 28.08.2025

# Motivation: Research context

## Cloud-native Database Management Systems:

- dynamic, multi-tenant environment
- resource utilization, fault tolerance
- flexibility

## General-purpose OSes

- heavyweight and resource-hungry

# Research gap

Unikernels - alternative to bulky OSes

- specialized and lightweight
- reduced attack surface
- single-address-space
- **BUT:** limited flexibility

How can we dynamically and securely reconfigure unikernel-based DBMS?

**Provide reconfiguration through a mechanism that enables safe and dynamic execution of extensions within the unikernel**

## **System design goals:**

- Live adaptability
- Lightweight-ness
- Performance
- Safety
- Portability

# Outline



- ~~Motivation~~
- Background
- Design
- Implementation
- Evaluation

## Extended Berkley Packet Filter (eBPF)

- OS level reconfiguration
- sandboxed programs in kernel space
- helper functions

## uBPF

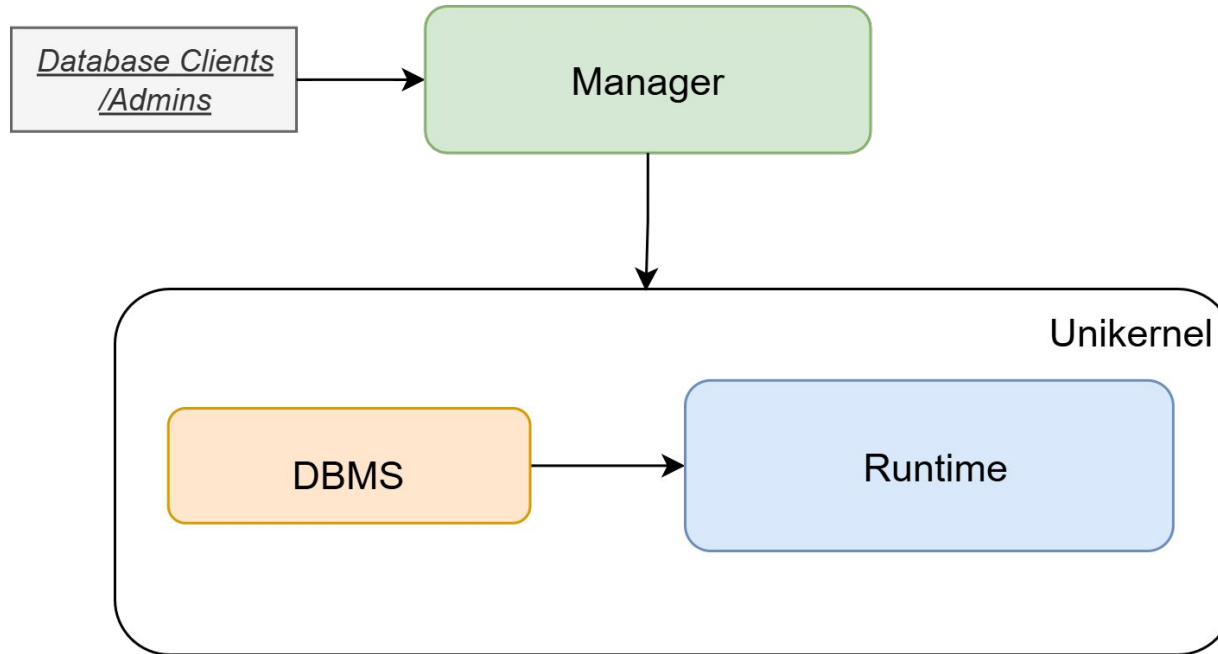
- user-space implementation of eBPF

# Outline

- ~~Motivation~~
- ~~Background~~
- Design
- Implementation
- Evaluation



# System overview



- DBMS
  - instrumented with hookpoints
- Runtime
  - execution engine for extensions
- Extensions
  - executed in an isolated environment
  - helper functions

The system is designed to evolve dynamically at runtime.

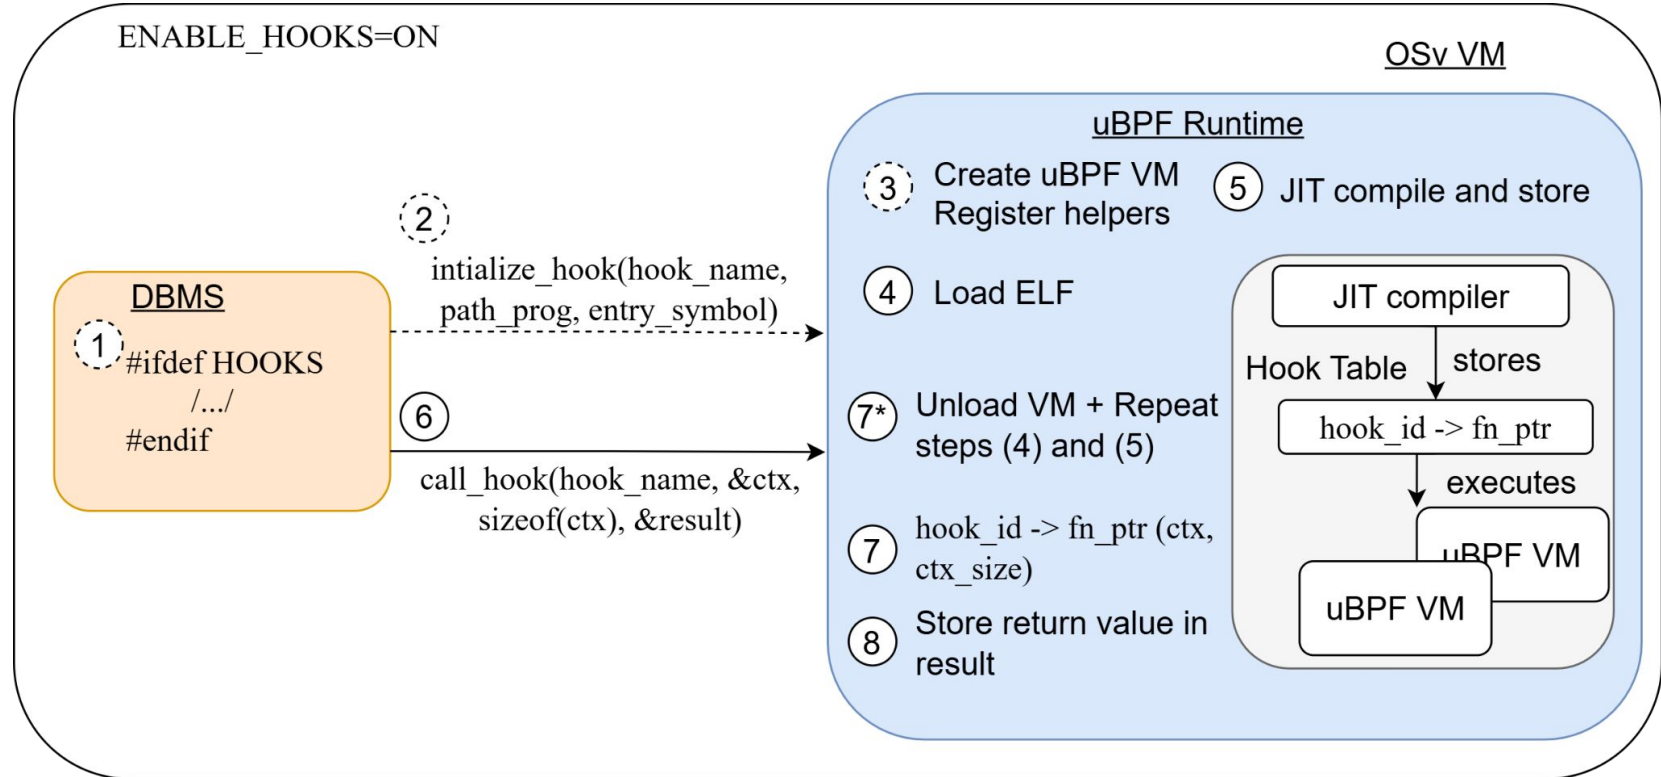
# Outline



- ~~Motivation~~
- ~~Background~~
- ~~Design~~
- Implementation
- Evaluation

- OSv unikernel - lightweight
- eBPF based extension mechanism using uBPF- safety and extensibility
  - helper functions
- JIT Compilation - performance
- DuckDB
  - instrumented with hookpoints

# Detailed Workflow



- Hookpoints
  - instrumented into DuckDB
  - replicate original behavior
- Helper Functions
  - internal functionality and state
  - e.g., scheduling queues, controlling buffer access

Hookpoint	func-	# Helpers
Hash	tion	1
ExecuteForever		9
ScheduleTask		1
Load		2
UpdateUsedMemory		3
AllocateNewBlock		2
Unpin		5

# Outline

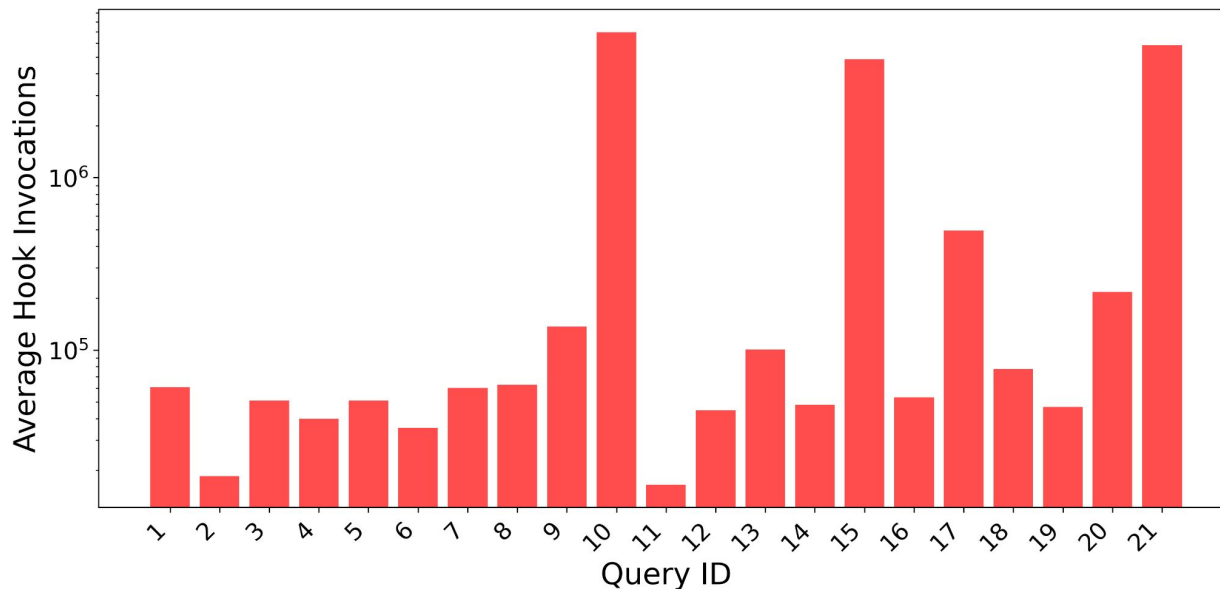
- ~~Motivation~~
- ~~Background~~
- ~~Design~~
- ~~Implementation~~
- Evaluation

- System under evaluation:
  - DuckDB instrumented with 7 hookpoints + uBPF Runtime
- Baseline:
  - DuckDB without hook support
- Research Questions:
  - Can the system preserve the performance of the DBMS?
  - Is the design lightweight, minimal memory overhead?
  - Does the mechanism improve reconfiguration times?




# Performance (Number of Hook Invocations)

TPC-H benchmark, 21 queries, scale factor 20



# Performance overhead


TPC-H benchmark, 21 queries, scale factor 20



The system introduces minimal performance impact

# Memory footprint

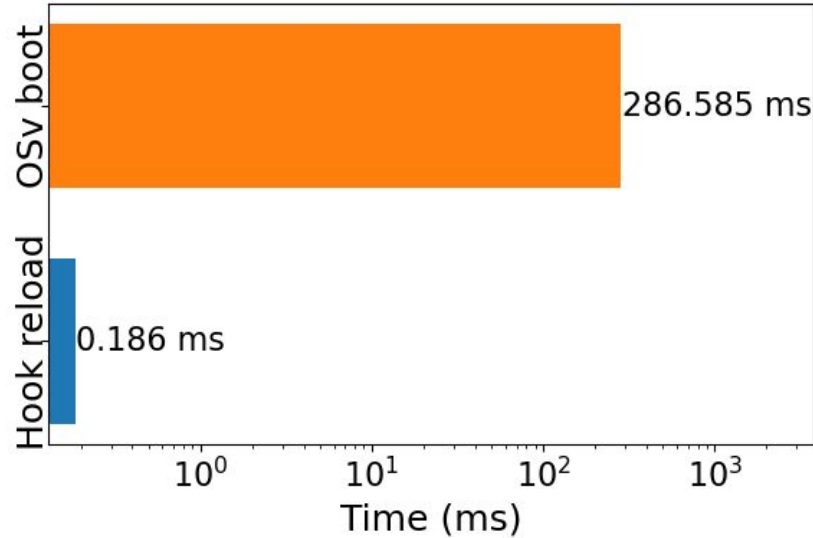
Compare OSv image library sizes



Insignificant memory overhead.

# Reconfiguration time

Reboot time vs Reload path



Live reconfiguration substantially outperforms static updates.

**Proposition:** Enable safe and dynamic reconfiguration at the DBMS/OS boundary through a lightweight extension mechanism

**Extensible, cloud-native DBMS deployments:**

- lightweight unikernel
- instrumented DBMS engine
- safe eBPF-based Runtime component