

# Airlift

## A Binary Lifter Based on a Machine-Readable Architecture Specification

Anders Choi

Advisor: Martin Fink

Chair of Computer Systems


<https://dse.in.tum.de/>




29.11.2024 – 30.05.2025

# Motivation: Binary Lifting

- Critical component in a wide range of tasks
- Lacks reuse for custom IRs
- Highly complex and error-prone
- Error-sensitive



Binary lifter implementation is **critical** and **non-trivial**!



**Cross-Architecture Lifter Synthesis** leverages an existing lifter for one ISA to **inductively generate** a lifter for another ISA

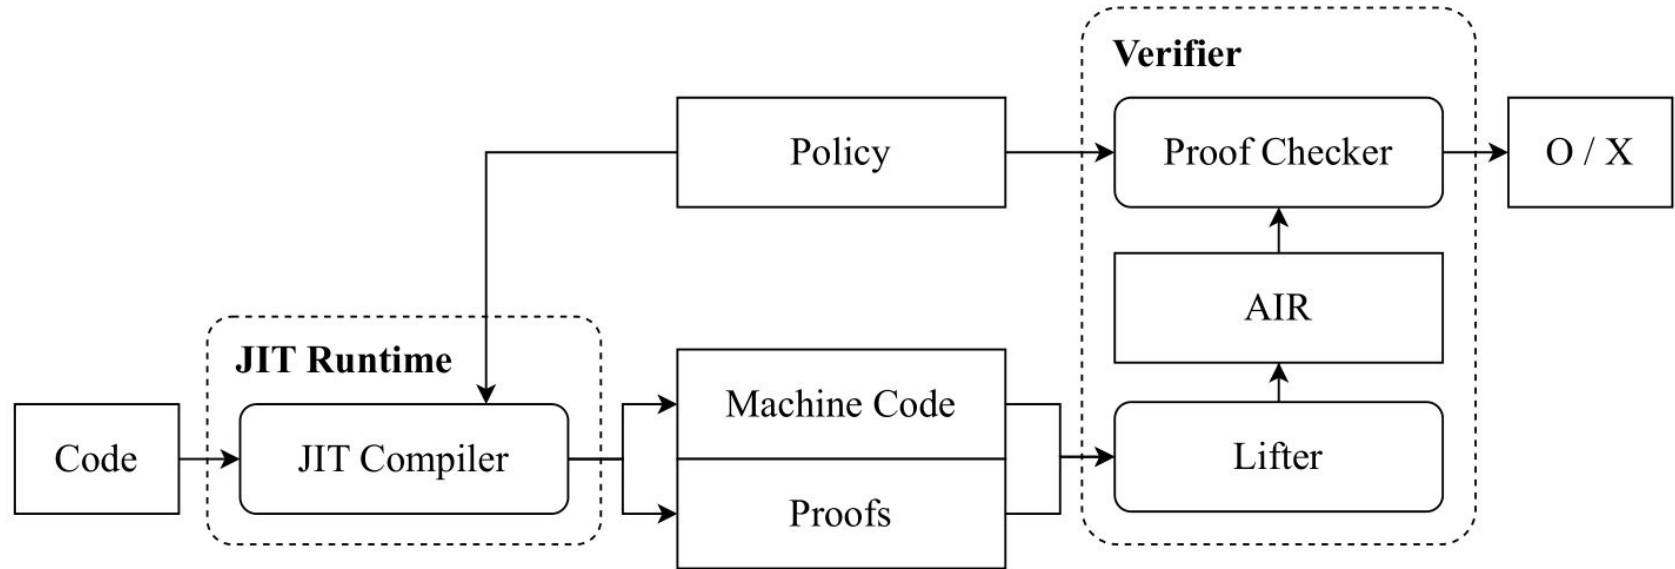
**Forklift** is a **neural lifter** that translates binary code to the target IR using a Transformer

**Lift-Offline** generates a lifter **from a formal specification**

SEFM 2018



# Context: TrustNoJit (TNJ)



A custom, lightweight IR for TNJ

- 3 data types
  - Unbounded integer
  - Fixed size integer
  - Boolean
- 37 instructions (vs. 100+ in LLVM IR)
- Block-based control flow

```
entry:
    v0 = i64.read_reg "x1"
    v1 = i64.read_reg "x2"
    v2 = i1.read_reg "n"
    v3 = bool.icmp.i1.eq v2, 0x1
    jumpif v3, block_0, block_1
block_0:
    jump block_2(v0)
block_1:
    jump block_2(v1)
block_2(v4: i64):
    write_reg.i64 v4, "x0"
```

How can we generate a **precise** binary lifter for a **security-critical** system?

## Challenges:

- Large volume of code
- Error-prone development
- Security-critical context

Transpile a **machine-readable architecture specification** into a binary lifter

## System design goals:

- High instruction coverage
- Semantic correctness
- ~~Performance and optimization~~

# ASL: Arm's Machine-Readable Architecture Specification



```
__decode A64
case (29 +: 3, 24 +: 5, 0 +: 24) of
...
  when (_, 'x101x', _) =>
    case (
      31 +: 1, 30 +: 1, 29 +: 1, 28 +: 1, 25 +: 3,
      21 +: 4, 16 +: 5, 10 +: 6, 0 +: 10
    ) of
      ...
        when (_, _, _, '1', _, '0100', _, _, _) => // csel
          __field sf 31 +: 1
          __field op 30 +: 1
          __field S 29 +: 1
          __field Rm 16 +: 5
          __field cond 12 +: 4
          __field op2 10 +: 2
          __field Rn 5 +: 5
          __field Rd 0 +: 5
          case (sf, op, S, op2) of
            when ('0', '0', '0', '00')
              => __encoding aarch64_integer_conditional_select
            ...
```

Instruction decode logic

Operand decode logic

Instruction Semantics

```
__instruction aarch64_integer_conditional_select
__encoding aarch64_integer_conditional_select
__instruction_set A64
__field sf 31 +: 1
__field op 30 +: 1
__field Rm 16 +: 5
__field cond 12 +: 4
__field o2 10 +: 1
__field Rn 5 +: 5
__field Rd 0 +: 5
__opcode 'xx011010,100xxxxx,xxxx0xxx,xxxxxxxx'
__guard TRUE
```

```
__decode
  integer d = UInt(Rd);
  integer n = UInt(Rn);
  integer m = UInt(Rm);
  integer datasize = if sf == '1' then 64 else 32;
  bits(4) condition = cond;
  boolean else_inv = (op == '1');
  boolean else_inc = (o2 == '1');
```

```
__execute
  bits(datasize) result;
  bits(datasize) operand1 = X[n];
  bits(datasize) operand2 = X[m];

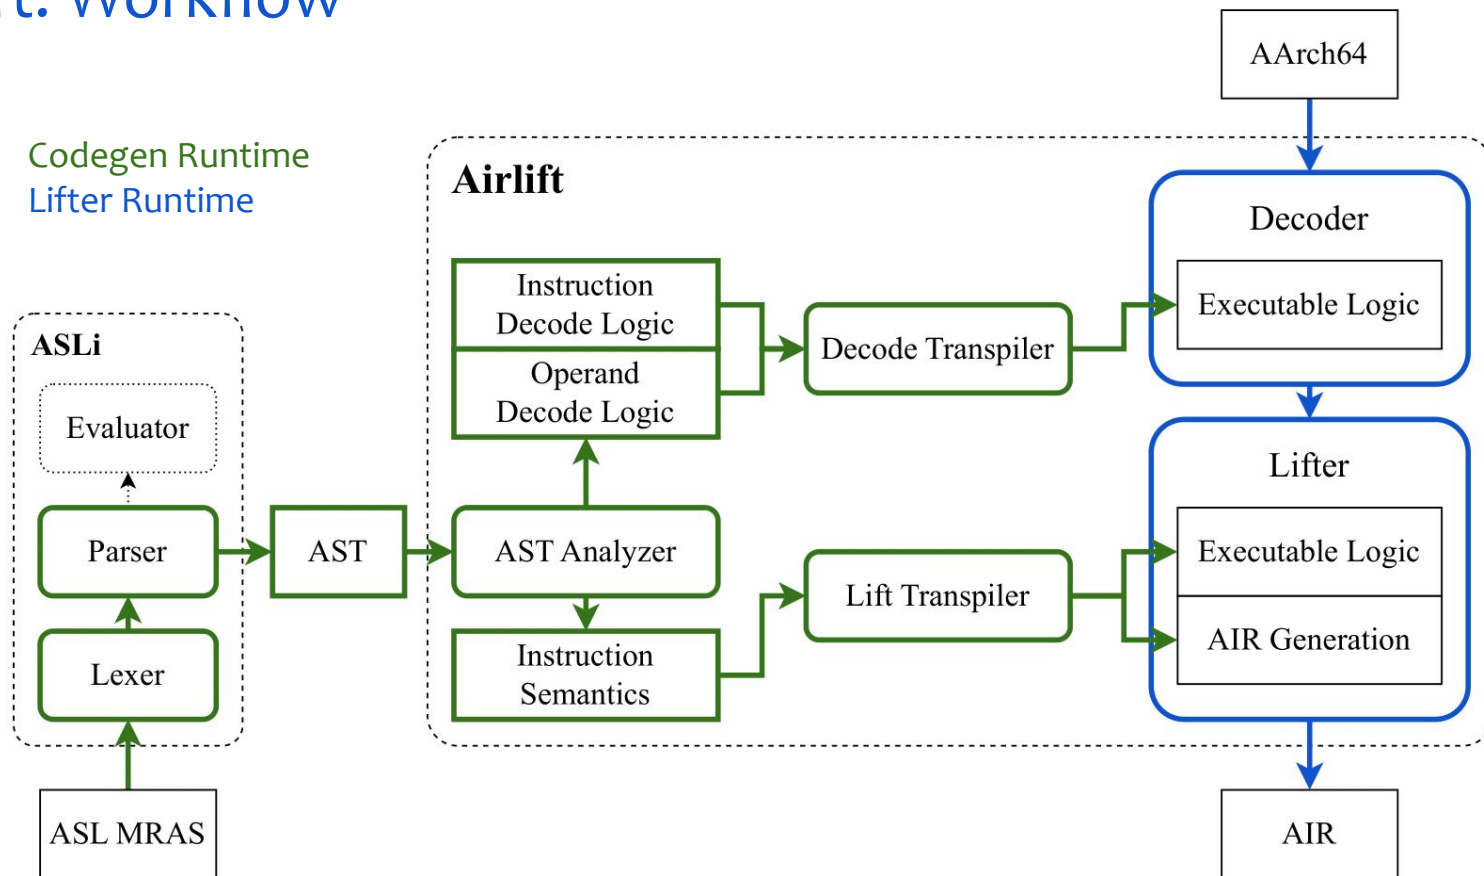
  if ConditionHolds(condition) then
    result = operand1;
  else
    result = operand2;
    if else_inv then result = NOT(result);
    if else_inc then result = result + 1;

  X[d] = result;
```

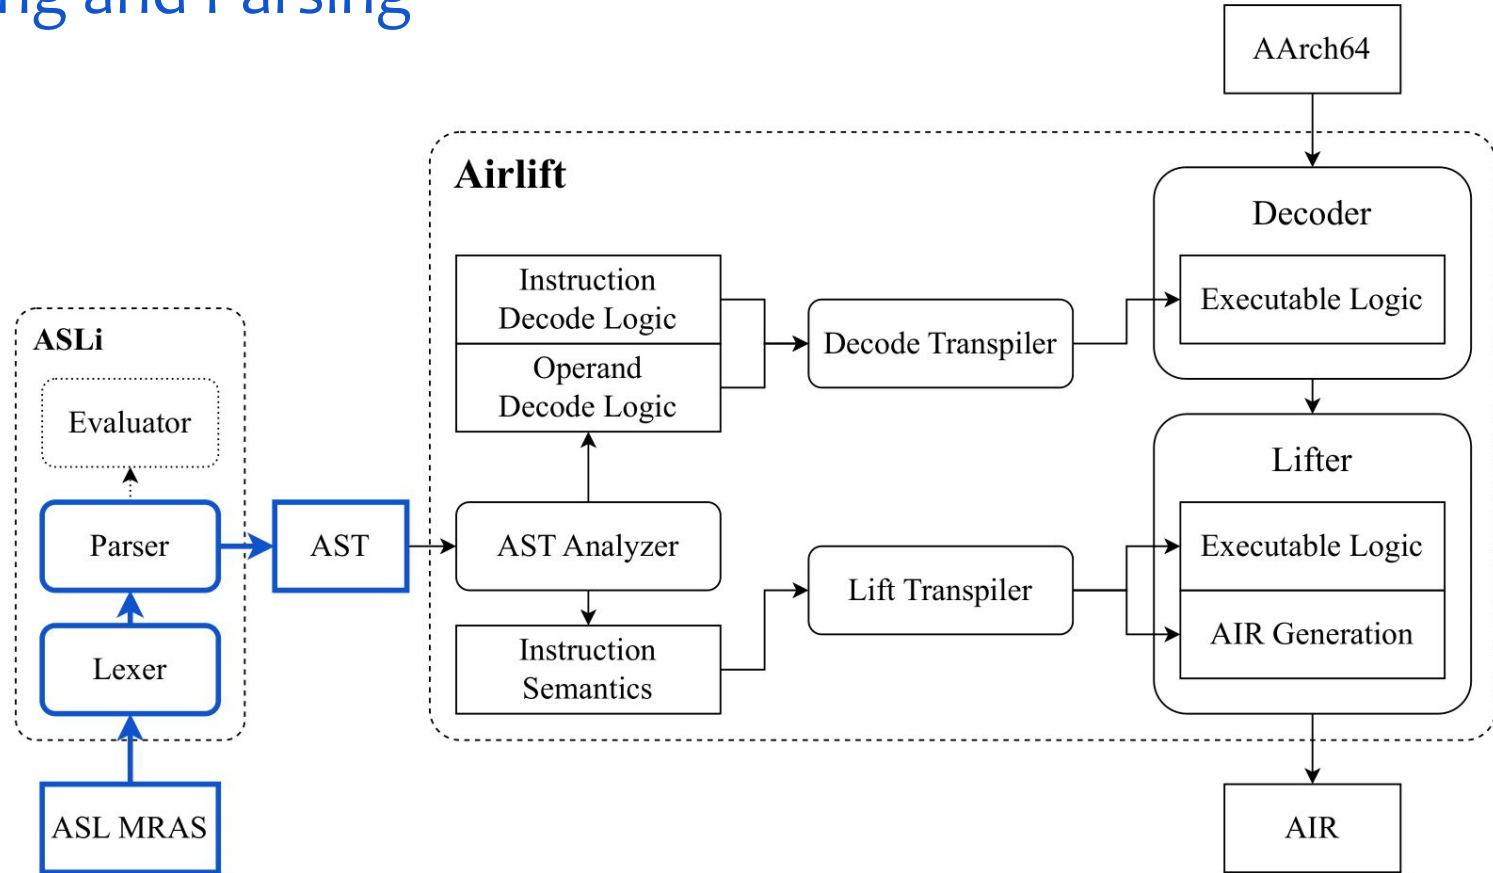


# Airlift: Workflow

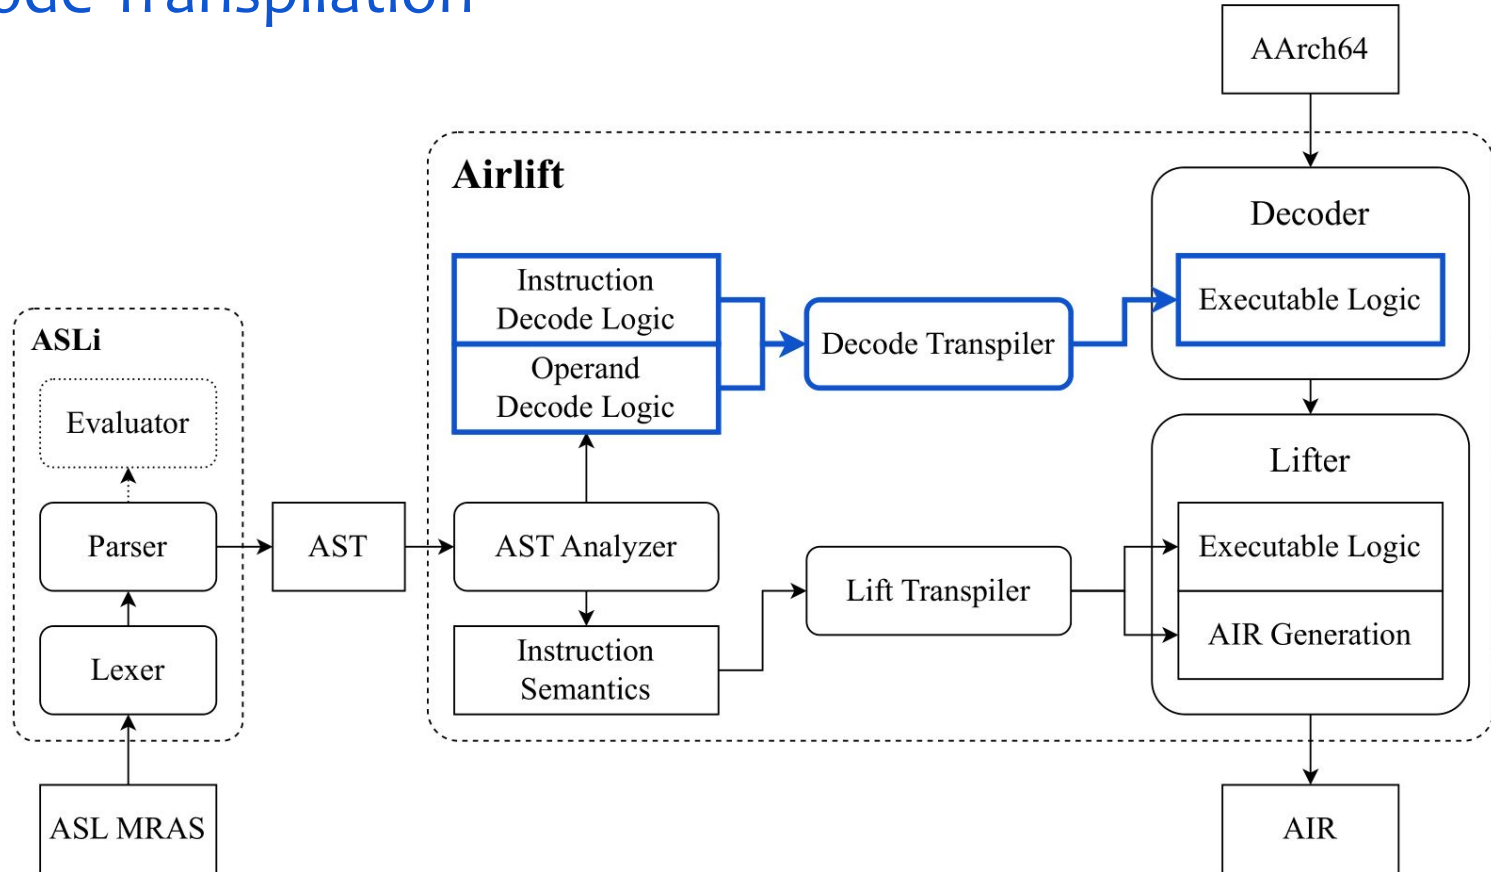
1. Codegen Runtime
2. Lifter Runtime



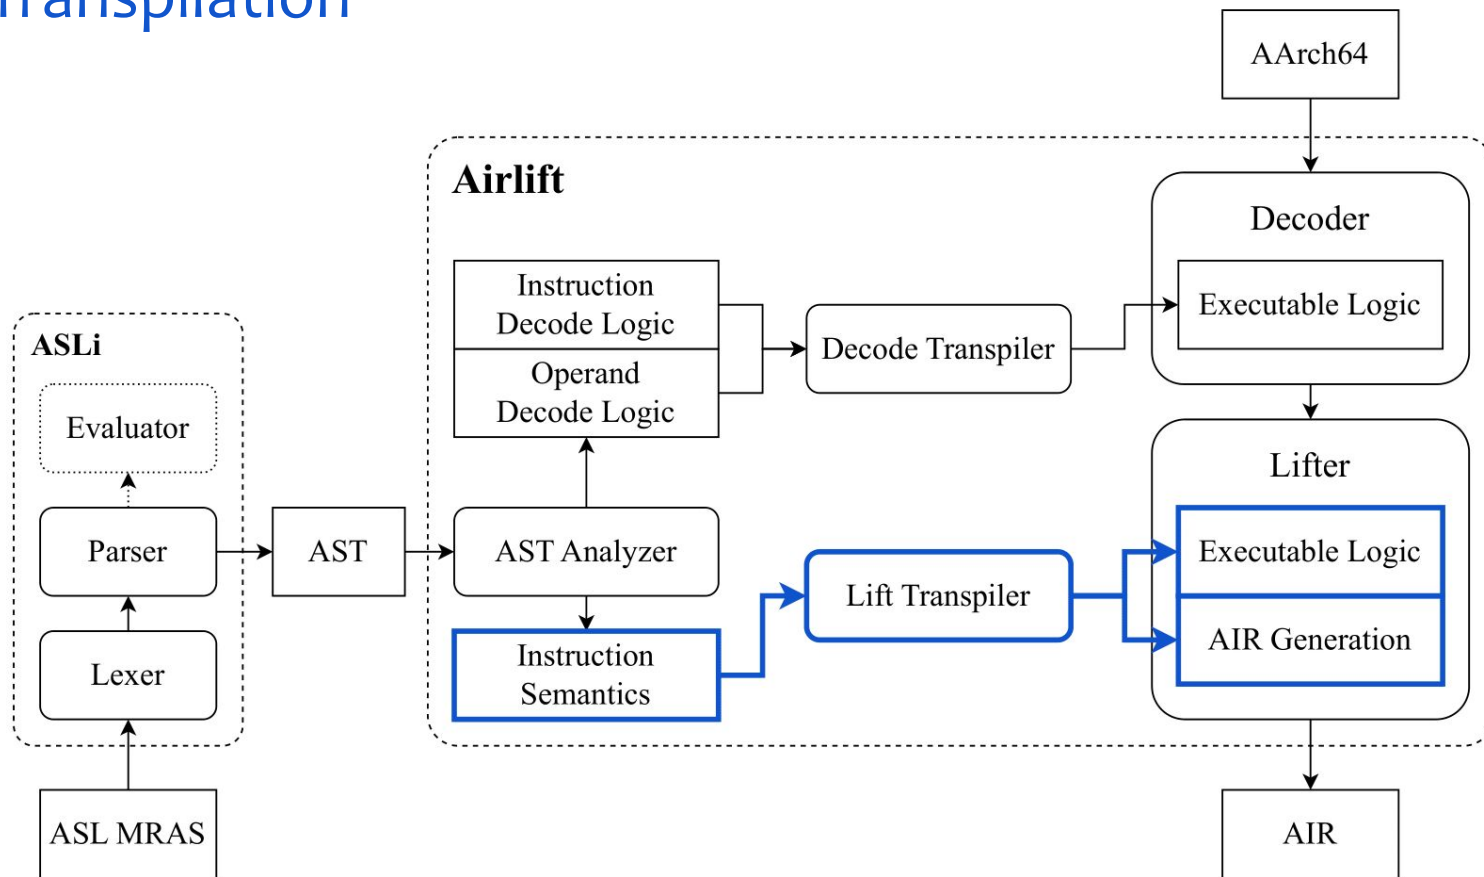
# Lexing and Parsing



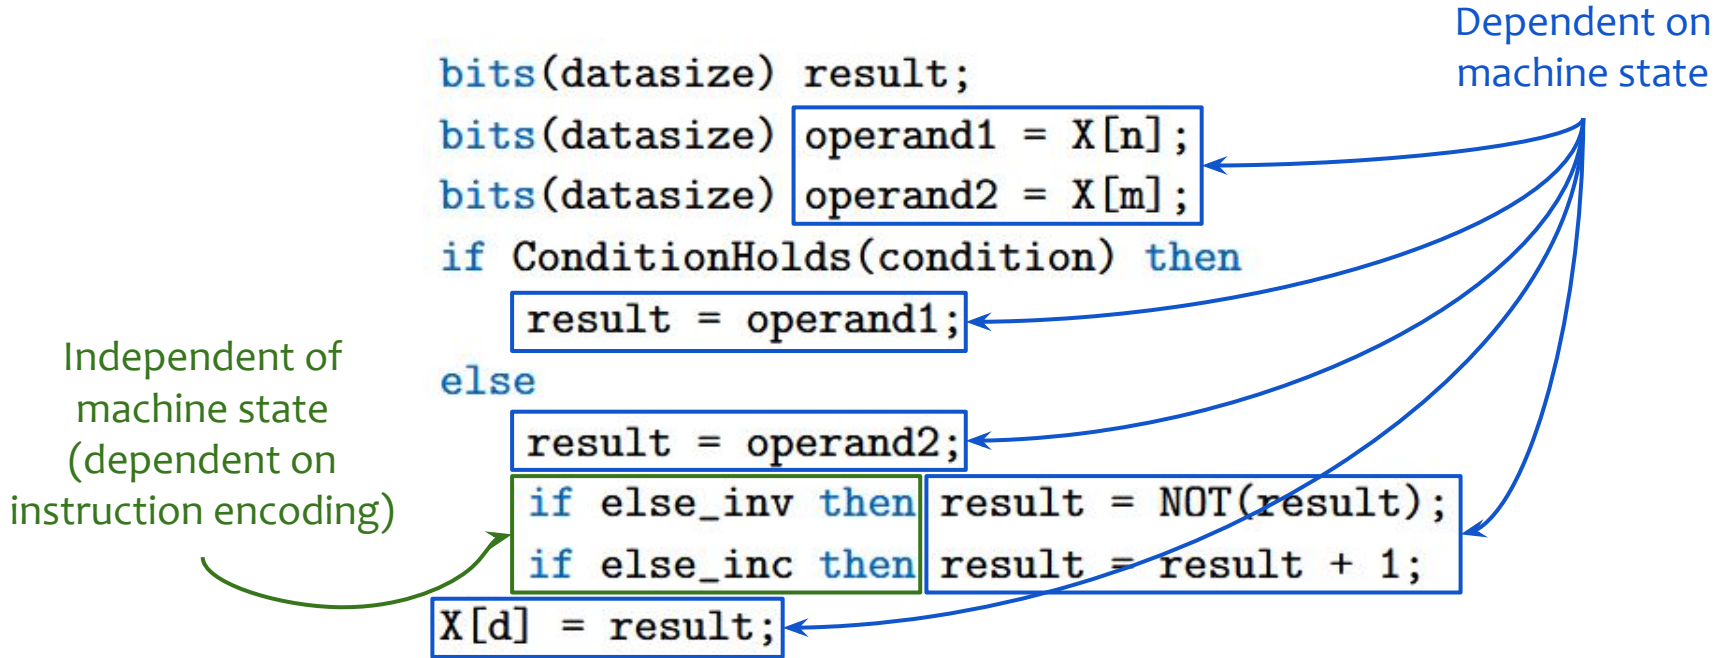
# Decode Transpilation



# Lift Transpilation



# Partial Evaluation Example



# Partial Evaluation Example

`--execute`

```
bits(datasize) result;  
bits(datasize) operand1 = X[n];  
bits(datasize) operand2 = X[m];
```

```
if ConditionHolds(condition) then
```

```
    result = operand1;
```

```
else
```

```
    result = operand2;
```

```
    if else_inv then result = NOT(result);
```

```
    if else_inc then result = result + 1;
```

```
X[d] = result;
```

`entry:`

```
    v0 = i64.read_reg "x1"
```

```
    v1 = i64.read_reg "x2"
```

```
    v2 = i1.read_reg "n"
```

```
    v3 = bool.icmp.i1.eq v2, 0x1
```

```
    jumpif v3, addr_0_block_0, addr_0_block_1
```

```
addr_0_block_0:
```

```
    jump addr_0_block_2(v0)
```

```
addr_0_block_1:
```

```
    jump addr_0_block_2(v1)
```

```
addr_0_block_2(v4: i64):
```

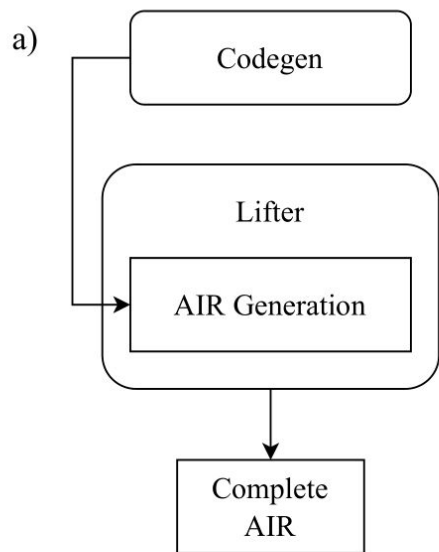
```
    write_reg.i64 v4, "x0"
```

✗

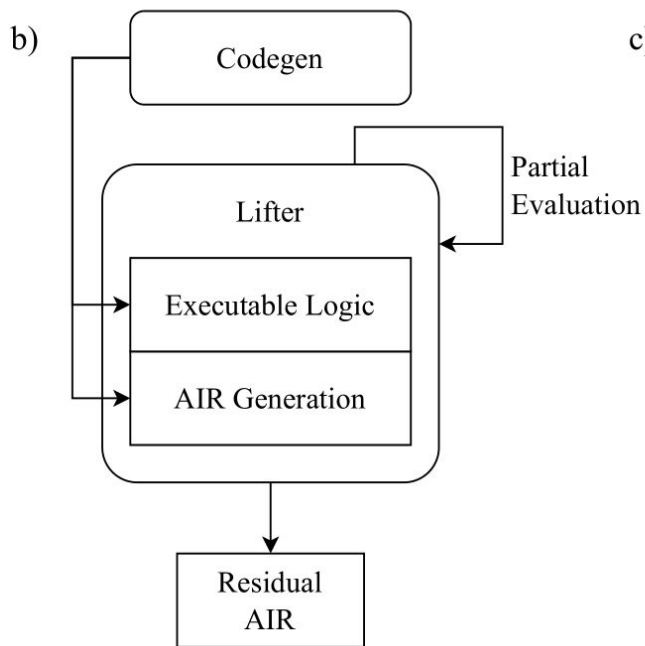
✗

# Partial Evaluation Alternatives

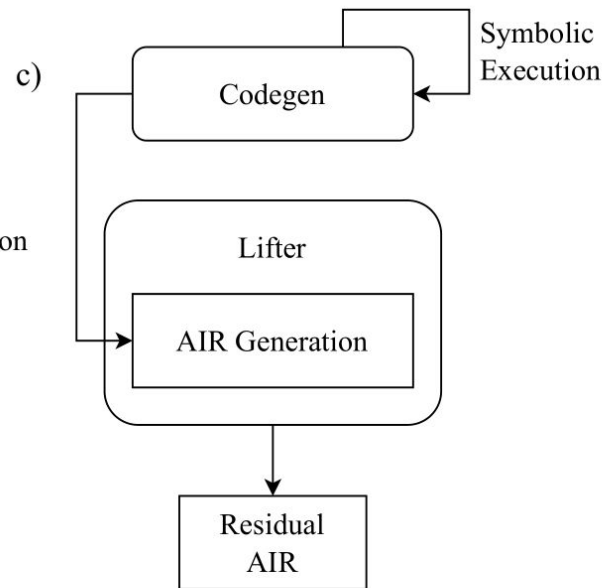
Naive



Airlift




Lift-Offline



# Implementation

- ASL interpreter (OCaml + Ott)
- AST serialized into custom JSON structure
- Airlift implemented in Rust
- FP/SIMD omitted (robust fallback generated)





**RQ1.** Can Airlift reliably decode and lift real-world programs?

**RQ2.** Is its performance on real-world programs sufficient for practical use?

**RQ3.** What are the time, instruction count, and block count characteristics across instruction types?

**RQ4.** Which instruction patterns cause inefficiencies, and what insights do they offer for improvement?

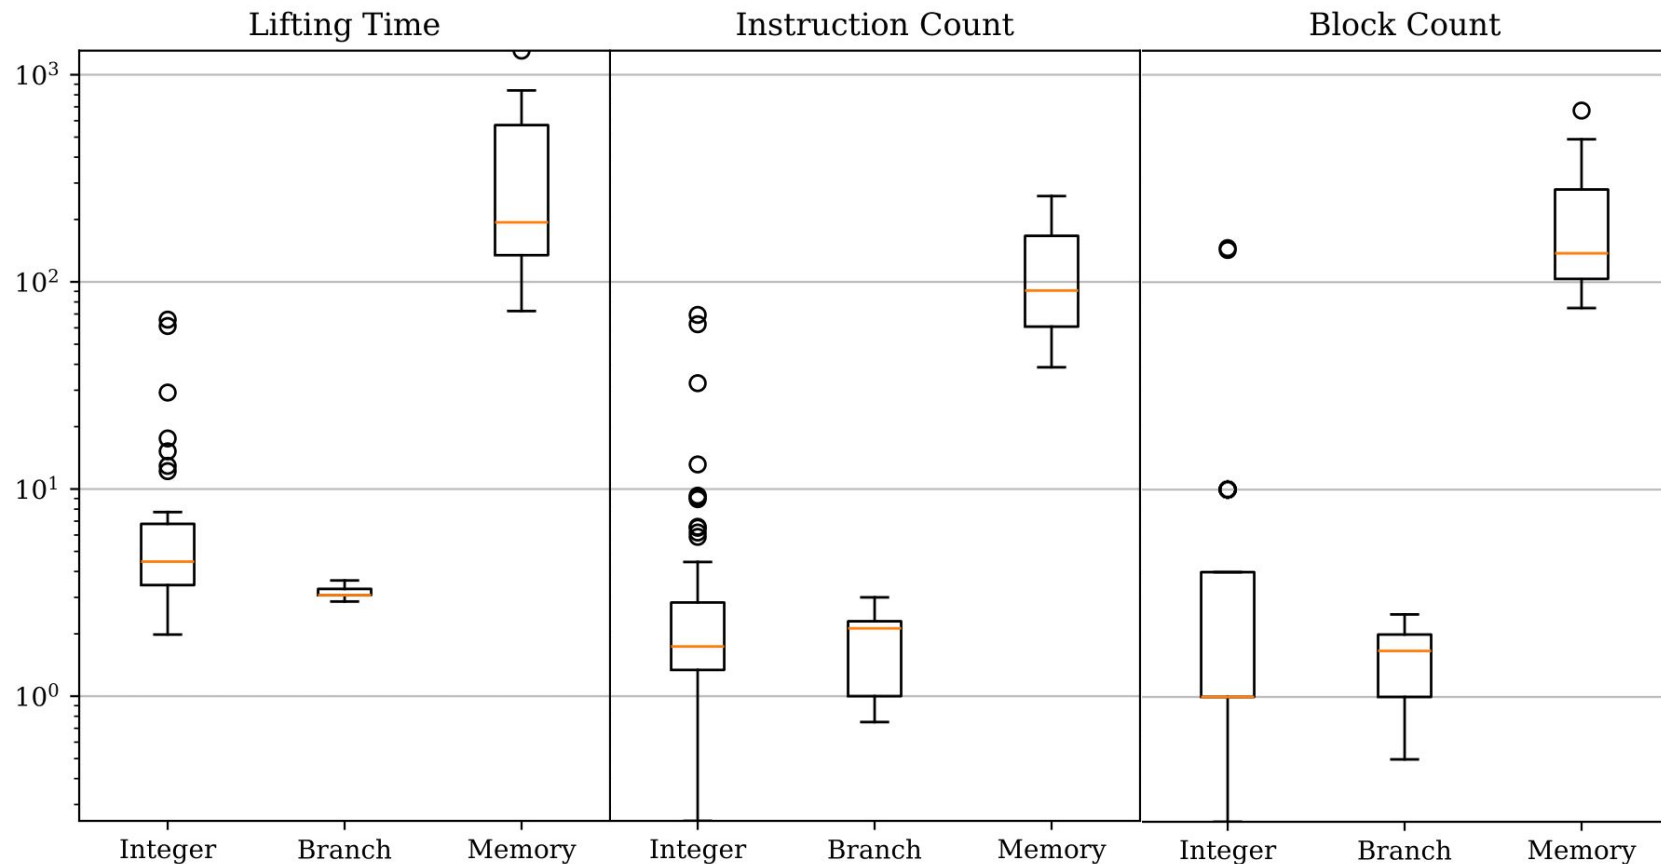
## Experimental setup

- CPU: AMD EPYC 7713P 64-Core Processor
- RAM: 991 GiB
- OS: NixOS

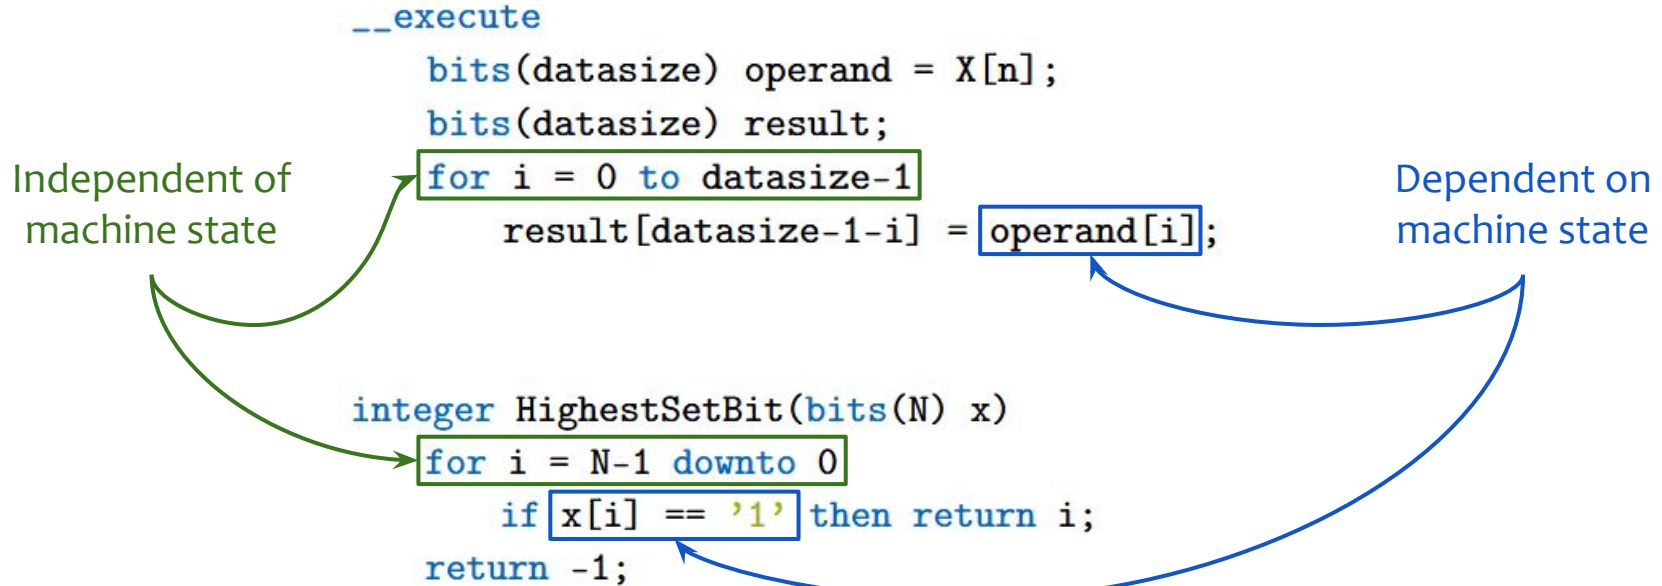
## Results

- Successfully decoded all instructions in 114 Sightglass WASM applications
- Successfully lifted all instructions except the deprioritized ones, notably FP/SIMD
- 1111 times slower than objdump

# Comparison with Manually Written Baseline Lifter



# Source of Inefficiency: Loops



## Airlift:

- Leverages Arm's machine-readable architecture specification
- Partially evaluates instruction semantics to reduce generated IR
- Focuses on correctness over performance

## Strengths:

- Strictly adheres to formal specification
- Reliably decodes and lifts real-world applications

## Limitation:

- Impractical lifting time → significant room for optimization

**Try it out!**

<https://github.com/TUM-DSE/airlift>