

# Disaggregated & Vendor-Generic APUs with CXL

Victor Fritz Trost

Advisor: Nicolò Carpentieri, Anatole Lefort

Chair of Computer Systems

<https://dse.in.tum.de/>



# Motivation: Memory bottleneck in GPUs



- **Increased VRAM demand**
  - AI models, data sets keep getting bigger, **can't always fit in VRAM**
- **Inefficient copy-then-execute model**
  - Requires workload-specific data splitting and partitioning
- **Inefficiency of Unified Virtual Memory (UVM):**
  - Assisted through software

Which path to a **more efficient UVM?**

# Motivation: APUs (Accelerated Processing Units)



## APU: CPU-GPU tightly integrated on a single die

- Provides a “**hardware UVM**”: efficient and convenient solution for GPGPU
  - Developed to make GPU memory more efficient
- Currently **limited to a single platform**, or single chips
  - AMD Instinct MI300a: Single chip with integrated memory
  - Nvidia GH200, NVLink C2C: Separate chips, but single (proprietary) platform
  - NVLink Fusion: Nvidia + multiple OEMs, but not yet available
- APU-like design cannot be applied to arbitrary CPU, GPU pair

Could we couple any CPU, GPU into a *logical APU*?

# Research Gap: CPU-GPU coherent memory with CXL



- CXL is promising for a generic APU design
  - Enables **cache-coherent** CPU-GPU memory sharing
  - **External protocol:** can connect discrete (off-the-shelf) CPU, GPU devices
  - Open standard

But:

- No existing CXL CPU-GPU coherent hardware
- No simulation models



We don't know what performance to expect from a CXL-based APU

# Research Statement



Can an APU-like CPU/GPU interface be efficient with CXL?

## We need to:

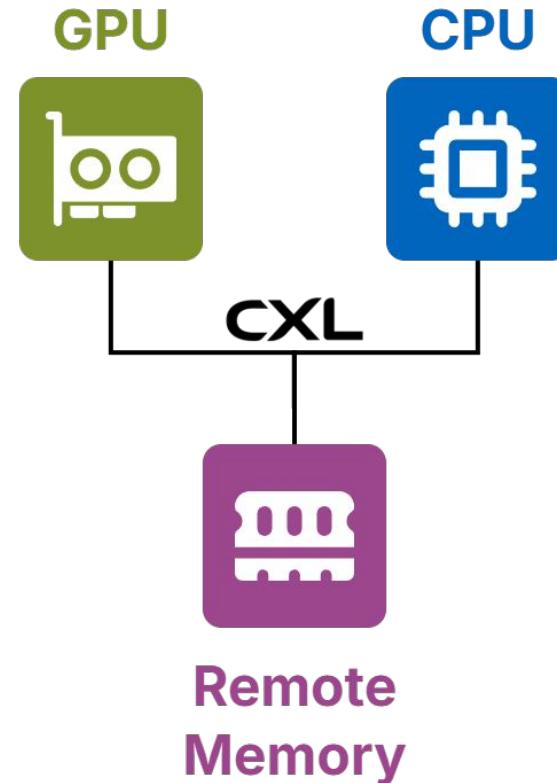
- Model performance of CXL-based APUs:
  - estimate **slowdowns of remote memory** for GPU kernels
  - estimate **gains with CPU-GPU cache coherence**

Thus we propose:

A simulation model for APUs with CXL memory

- **Goal #1 Programmability:** Single address space
- **Goal #2 Efficiency:** Hardware cache coherence
- **Goal #3 Genericity:** Platform agnostic

- **Efficient UVM:**  
hardware based
- **Unified Memory Pool:**  
scalable beyond VRAM limits
- **Platform independent:**  
not locked into a specific vendor
- **CXL-based**  
disaggregation-ready



# Background

State-of-the-Art GPU simulations:

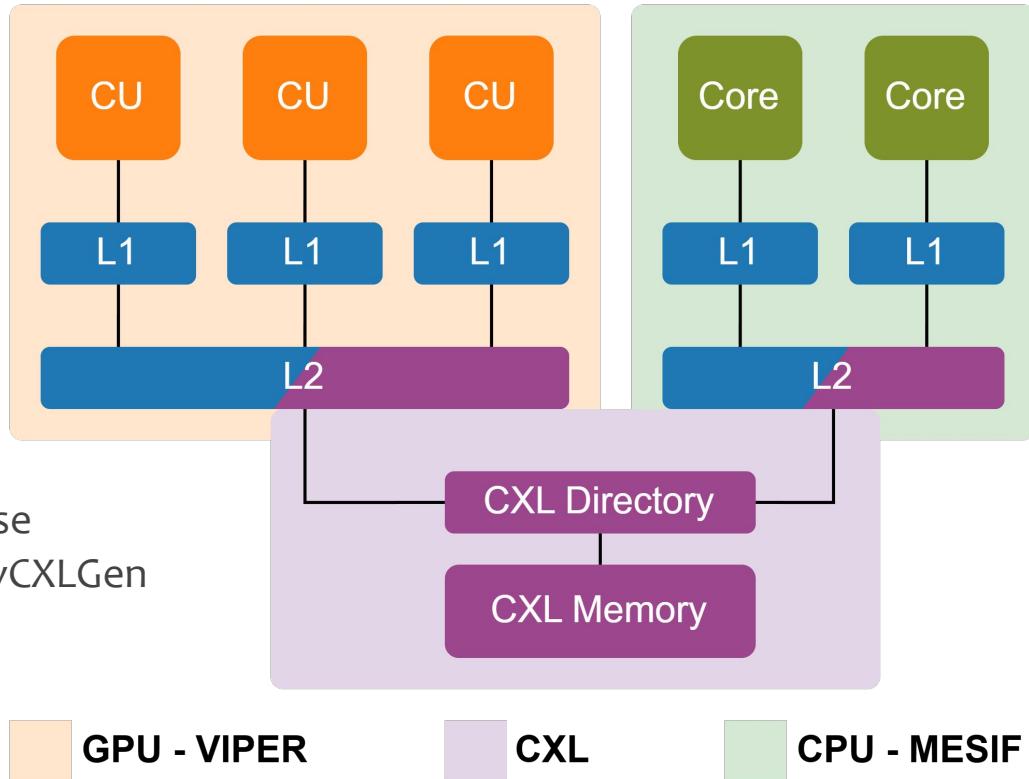
- **GPGPU-Sim:**
  - Simulates NVIDIA GPUs, limited CPU/UMA simulation, focus only on GPU kernels
- **Accel-Sim:**
  - Based on GPGPU-Sim, but trace-based, similar drawbacks
- **gem5's GPU models:**
  - Simulates AMD GCN, out-of-the-box support for APU



Quick proof-of-concept by extending gem5's APU model

# Implementation - Overview

- **CPU-side: MESIF protocol**
  - Generated using Synapse
- **GPU-side: VIPER protocol**
  - Already in gem5, extended
- Both protocols **bridged** to CXL
  - CXL domain generated with Synapse
  - MESIF-CXL bridge generated with vCXLGen
  - VIPER-CXL bridge hand-written

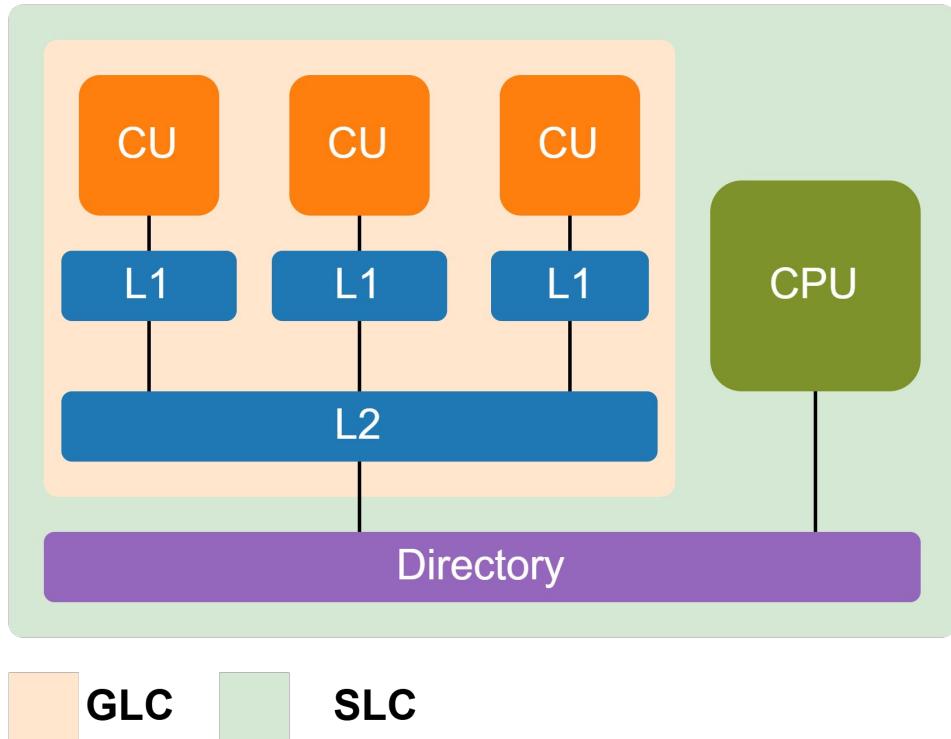


## Why MESIF?

- **Forwarding (F):** cache-to-cache sharing to minimize latency
- Intel uses MESIF-like protocol in multi-socket (NUMA) servers
- We use **directory-based** approach
  - Simpler to generate bridge, better scalability
- Gem5 doesn't have MESIF
  - We write MESIF specifications and generate it with Synapse PCC-to-SLICC compiler

# GPU Implementation - VIPER

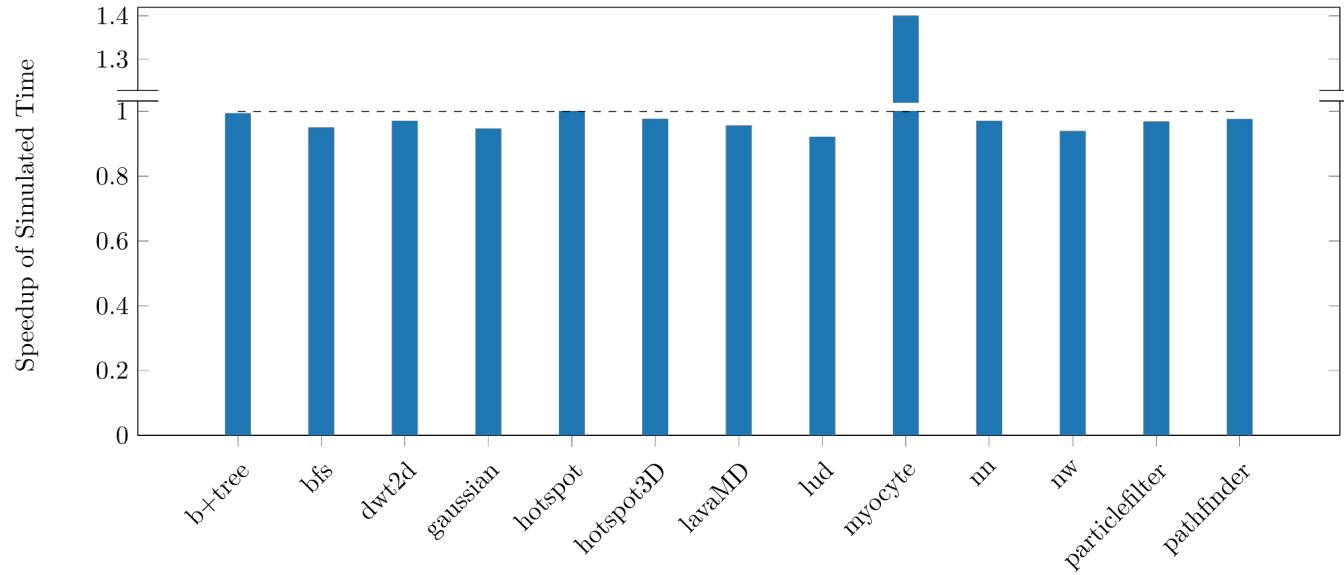
- VIPER provided by gem5
- **Adapt** only **L2** to CXL-bridge
- CXL bridge implementation:  
VIPER is a scoped protocol,  
L2 data consistency changes wrt.  
access scope:
  - **SLC**: system-wide (incl. CPU)
  - **GLC**: only GPU-wide
- Uses **SLC** and **GLC** scopes
  - Introduce a new incoherent P state in CXL to model GPU-wide coherent data (GLC)



## How does the CXL-based disaggregated APU compare to the integrated APU chip?

- Use Rodinia benchmark suite
  - Heterogeneous benchmark suite focusing on GPUs
  - Cross platform and domain
  - Implemented in CUDA
  - We translate to HIP (AMD GPGPU stack) with the *hipify* tool
- Use original gem5 model as baseline (AMD VIPER APU)
  - Original GPU latency as CXL latency (60ns)
  - Original model uses local DRAM -> faster CPU accesses than with CXL memory
- 4 CPU cores (MESIF), 4 GPU CUs (VIPER)

# Evaluation - Overall

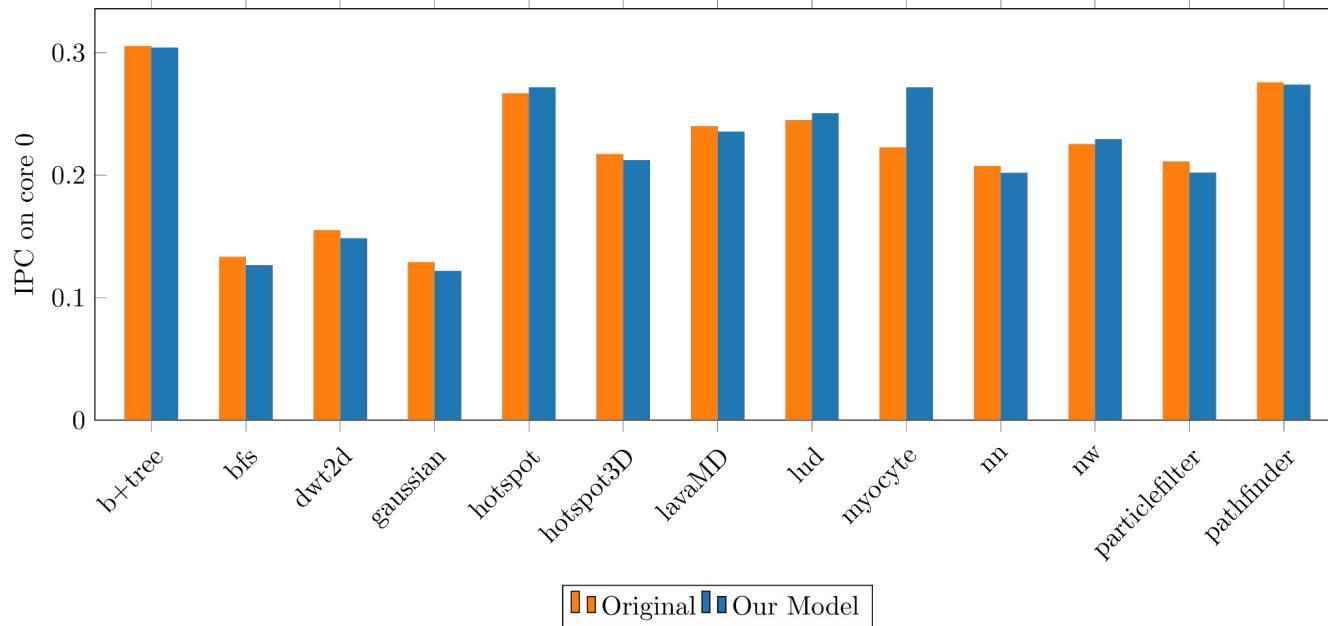


At similar latencies to original runtime



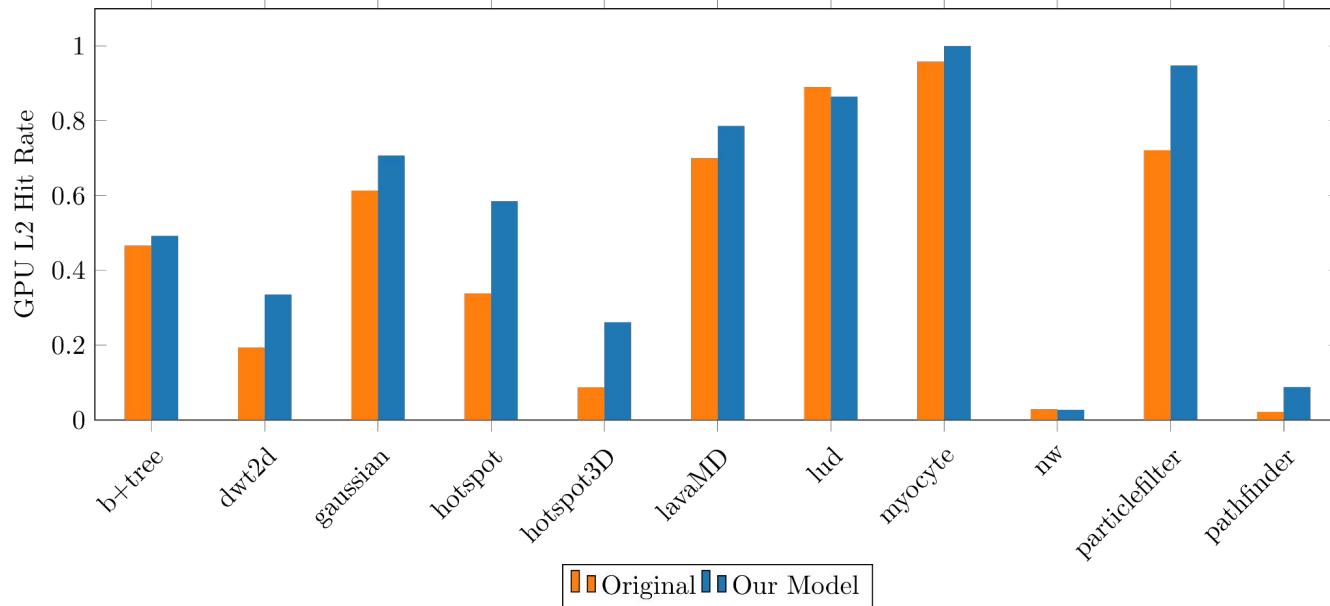
Comparable performance to original VIPER implementation

# Evaluation - CPU



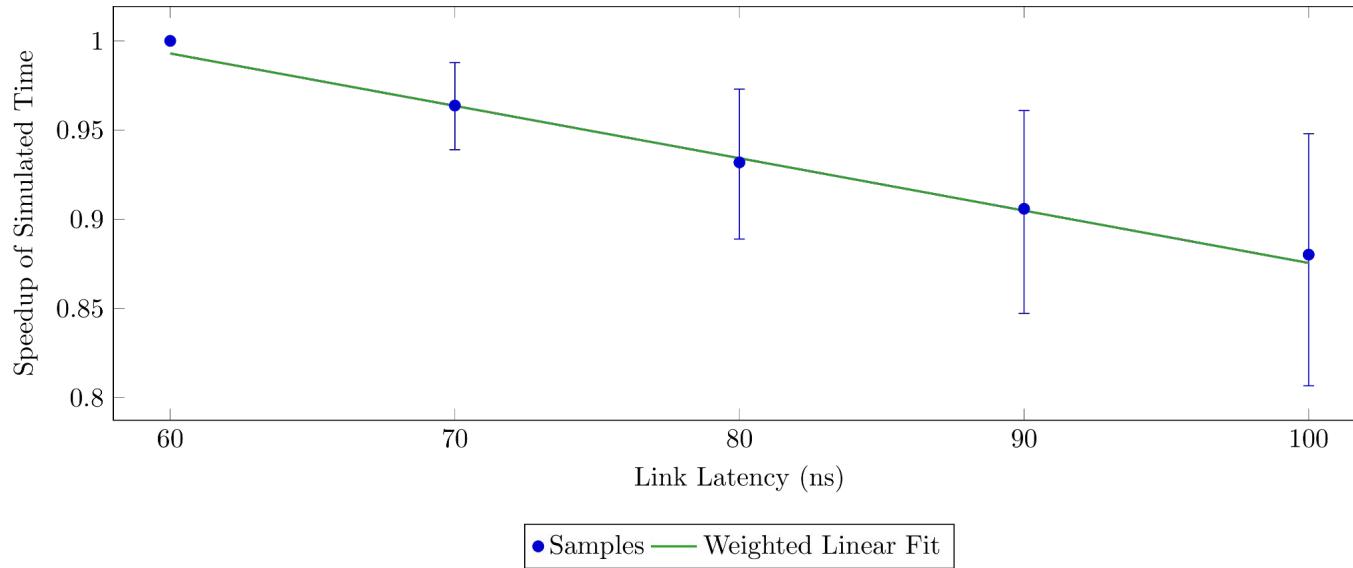
- Slightly worse IPC w/ CXL
- Due to higher memory latency with CXL (from CPU), original uses a different topology
- Original “bridges” GPU’s L2 directly to CPU’s LLC

# Evaluation - GPU



- Overall higher hit rate
- Expected due to write-through nature of original VIPER

# Evaluation - Link Latency

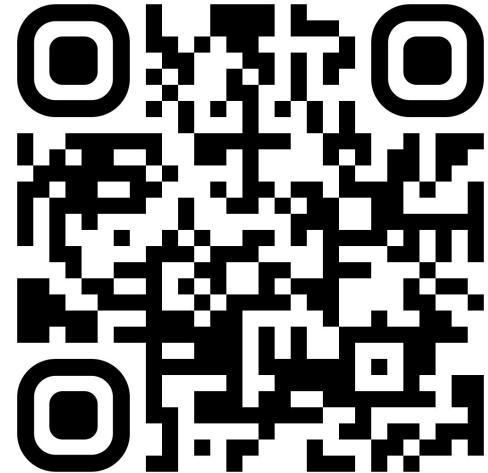


- Linear slowdown with increasing CXL link latency
- Performance degradation comprised between 5-20% for 1.5x link latency

# Conclusion

We achieve:

- A generic APU using CXL shared memory
- With well-scaling performance, similar to gem5's original APU



Try it!

[github.com/TrostV/cxl-apu](https://github.com/TrostV/cxl-apu)



CXL-based APUs are as efficient as integrated APUs



CXL emerges as a valuable option in the VRAM limited AI-era.